

Advanced Topics in Deep Learning

MINI-PROJECT 02

**DEEP REINFORCEMENT LEARNING
DQN AND PPO IN LUNARLANDER V3**

Table of Contents

1. Introduction	4
2. Environment Analysis.....	4
2.1 State Space.....	4
2.2 Action Space	4
2.3 Reward Function	5
2.4 Termination Rules.....	5
3. Experimental Setup	6
3.1 Total Environment Steps	6
3.2 Hyperparameter Configurations	6
3.3 Random Seeds	7
3.4 Evaluation Protocol	7
3.5 Best-Model Selection.....	7
3.6 Reproducibility	7
4. Evaluation.....	20
4.1 Per-Algorithm Results.....	21
4.2 Cross-Algorithm Comparison	24
4.3 Statistical Significance	26
4.4 Baseline Comparison	26
4.5 Agent Behavior Analysis.....	27
4.6 Sample Efficiency Analysis	28
4.7 Training Milestone Timeline	29
4.8 Final vs Best Model Comparison	30
4.9 Hyperparameter Exploration Summary	31
4.10 Visualizations	32
5. Experimental Setup	34
6. Conceptual Discussion.....	35
6.1 Value-Based vs Policy-Based Learning	35
Value-Based Learning	35
Policy-Based Learning	35
6.2 Overestimation and Instability in Value-Based Algorithms	36
Overestimation Bias	36
Structural Instability: The Deadly Triad.....	36
6.3 PPO Clipping and Generalised Advantage Estimation (GAE)	36
PPO Clipped Objective	36
Generalised Advantage Estimation (GAE).....	37

6.4 Epsilon-Greedy vs Entropy-Driven Exploration.....	37
Epsilon-Greedy Exploration (DQN)	37
Entropy-Regularised Exploration (PPO).....	38
6.5 Conceptual Summary	38

1. Introduction

Deep Reinforcement Learning (DRL) has become a central paradigm for solving complex sequential decision-making problems, particularly those involving continuous state spaces, stochastic dynamics, and delayed rewards. Among the wide range of DRL algorithms, Deep Q-Network (DQN) and Proximal Policy Optimization (PPO) represent two fundamentally different approaches: value-based learning and policy-based learning. Comparing these two algorithms provides valuable insight into how different learning mechanisms behave under the same environment conditions.

In this project, we investigate the performance of DQN and PPO in the LunarLander-v3 environment from Gymnasium. This environment simulates the control of a lunar lander that must descend and touch down smoothly within a designated landing zone. The agent receives an 8-dimensional continuous state vector and must choose among four discrete actions, making the task suitable for both value-based and policy-based methods. The environment includes dense rewards, penalties for fuel consumption, and strong negative rewards for crashes, creating a challenging control problem that requires stability, precision, and efficient exploration.

The goal of this work is to train, evaluate, and compare DQN and PPO under controlled experimental conditions. We analyse their learning curves, sample efficiency, stability across multiple random seeds, and qualitative behaviour through recorded simulations. Additionally, we examine theoretical aspects such as overestimation bias in DQN, the role of clipping and Generalized Advantage Estimation (GAE) in PPO, and the differences between epsilon-greedy and entropy-driven exploration.

By combining quantitative metrics with qualitative observations, this study aims to provide a comprehensive understanding of how DQN and PPO behave in the LunarLander-v3 environment, highlighting their strengths, limitations, and the impact of hyperparameter choices on the final performance.

2. Environment Analysis

2.1 State Space

The LunarLander-v3 environment provides an 8-dimensional continuous state vector describing the physical configuration of the lander at each timestep: 1. Horizontal position (x) relative to the landing zone 2. Vertical position (y) above the ground 3. Horizontal velocity (vx) 4. Vertical velocity (vy) 5. Lander angle (θ) 6. Angular velocity (ω) 7. Left leg contact indicator (0 or 1) 8. Right leg contact indicator (0 or 1)

These variables allow the agent to infer its location, orientation, stability, and motion, which are essential for controlling the descent.

2.2 Action Space

The action space is discrete with four possible actions:

- 0 - Do nothing

- 1 - Fire left thruster (pushes the lander to the right)
- 2 - Fire main thruster (reduces vertical speed)
- 3 - Fire right thruster (pushes the lander to the left)

These actions allow the agent to control both horizontal movement and vertical descent.

2.3 Reward Function

The reward function is dense and designed to encourage smooth and stable landings while penalizing unsafe or inefficient behaviour.

Positive rewards include: - Moving closer to the landing zone - Reducing horizontal and vertical velocity - Maintaining a stable angle - Each leg touching the ground (+10 each) - Successful landing (+100 to +140)

Penalties include: - High velocities - Large tilt angles - Excessive use of the main thruster (fuel cost) - Crashes (around -100)

A score of **200 or higher** typically indicates a successful landing.

2.4 Termination Rules

An episode terminates when one of the following occurs: 1. **Successful landing** within the designated zone 2. **Crash** due to excessive speed or unstable angle 3. **Leaving the screen boundaries** 4. **Reaching the maximum number of steps** allowed by the environment (1000 timesteps)

These termination conditions apply equally to both PPO and DQN agents.

```
Session prefix: dqn_ppo
Algorithms: ['dqn', 'ppo']
Seeds: [42, 123, 3407]
Wind enabled: False
Total timesteps per seed: 1,500,000
Evaluation episodes per seed: 20
Chart update frequency: every 10 episodes
Checkpoint frequency: every 100 episodes
Eval callback frequency: every 25,000 timesteps (20 episodes)
Solved threshold: 200
Device: cpu
```

3. Experimental Setup

This section describes the training and evaluation methodology used to compare DQN and PPO in the LunarLander-v3 environment. All experiments were conducted under controlled and reproducible conditions to ensure a fair comparison between the two algorithms.

3.1 Total Environment Steps

Both DQN and PPO were trained using a total of **1,500,000 environment steps per seed**. Since the project uses 3 seeds, this results in:

- 3 seeds x 1,500,000 steps = **4,500,000 training steps per algorithm**

Using the same number of interactions ensures that the comparison between DQN and PPO is based on equal sample budgets.

3.2 Hyperparameter Configurations

The hyperparameters for each algorithm were selected based on Stable-Baselines3 recommendations and refined through empirical testing. Tables 1 and 2 summarize the final configurations used for training.

Hyperparameter	Value
Policy	MlpPolicy
Learning rate	linear_schedule(6.3e-4)
Buffer size	750,000
Batch size	128
Gamma (discount factor)	0.99
Learning starts	50,000
Exploration strategy	ϵ -greedy
Exploration fraction	0.12
Final ϵ	0.1
Target update interval	250
Train frequency	4
Gradient steps	4
Network architecture	256 x 256 MLP

Table 1 - DQN Hyperparameters

Hyperparameter	Value
Policy	MlpPolicy
Learning rate	0.00025
n_steps	2048

Hyperparameter	Value
Batch size	64
n_epochs	10
Gamma (discount factor)	0.999
GAE λ	0.95
Clip range	0.2
Entropy coefficient	0.01
Network architecture	256 x 256 MLP

Table 2 - PPO Hyperparameters

These hyperparameters were kept constant across all seeds to isolate the effect of stochasticity in training.

3.3 Random Seeds

To evaluate robustness and training stability, both algorithms were trained using **3 different random seeds**: - Seed 1: **42** - Seed 2: **123** - Seed 3: **3407**

Each seed corresponds to a full independent training run of 1,500,000 steps.

3.4 Evaluation Protocol

Each trained model was evaluated over **20 deterministic episodes per seed**, for a total of 60 evaluation episodes per algorithm. Deterministic evaluation ensures that the policy is tested without exploration noise, providing a fair assessment of the learned behaviour.

3.5 Best-Model Selection

During training, an evaluation callback assessed the model every 25,000 environment steps using 20 deterministic episodes. The best model was selected based on a combined metric: **mean reward - standard deviation**. This metric favours models that are both high-performing and consistent, rather than models that achieve high mean reward with large variance.

A two-tier selection gate ensures that once any model crosses a mean reward of 200, only solved models can replace the current best.

3.6 Reproducibility

All experiments used fixed random seeds for Python, NumPy, PyTorch, and the Gymnasium environment. Deterministic algorithms were enabled in PyTorch to minimize non-determinism from GPU operations.

Python: 3.12.3
PyTorch: 2.10.0+cu130
Device: cpu

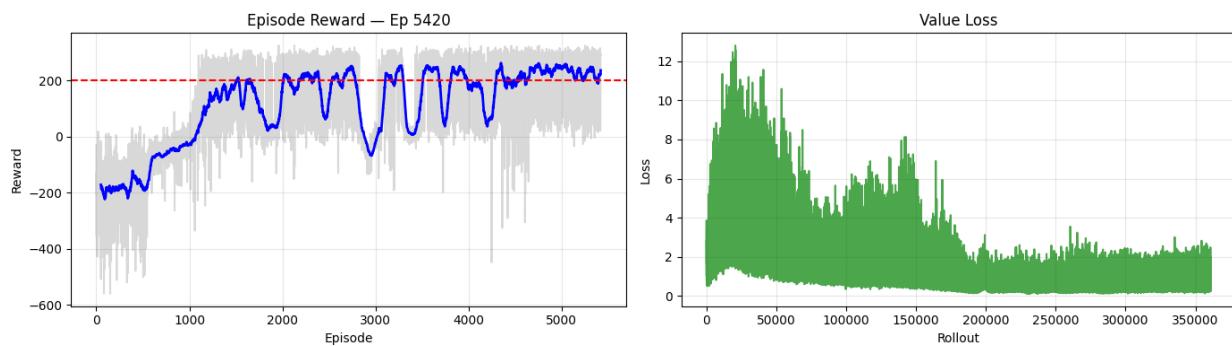
```
CUDA: 13.0
```

```
Observation space: Box([-2.5, -10., -10., -6.28318, 5.5, -10., -0., -0., 2.5, 2.5, 10., 10., 6.283185, 1., 1.], [2.5, 10., 10., 6.283185, 10., 10., 1., 1.], (8,), float32)
Action space: Discrete(4)
Initial observation: [-0.00170641, 1.4186276, -0.17285168, 0.3425566, 0.00198406, 0.03915349, 0., 0.]
```

```
=====
DQN | Seed 42
=====
```

Output()

```
'Episode 5420 | Last 50 Ep - Mean: 236.2 | Std: 95.5 | Min: 15.8 | Max: 315.9 | Success: 82%'
```



```
'[Checkpoint] Episode 5400 saved'
```

```
'Eval @ 1500000 steps | Reward: 276.50 +/- 16.66 [SOLVED] | Success: 100% | Score (mean-std): 259.85 | Best: 266.35'
```

```
Output()
```

```
'Episode 7790 | Last 50 Ep - Mean: -708.5 | Std: 349.2 | Min: -2311.8 | Max: -26
2.3 | Success: 0%'
```

Training time: 47.6 min (2856 s)

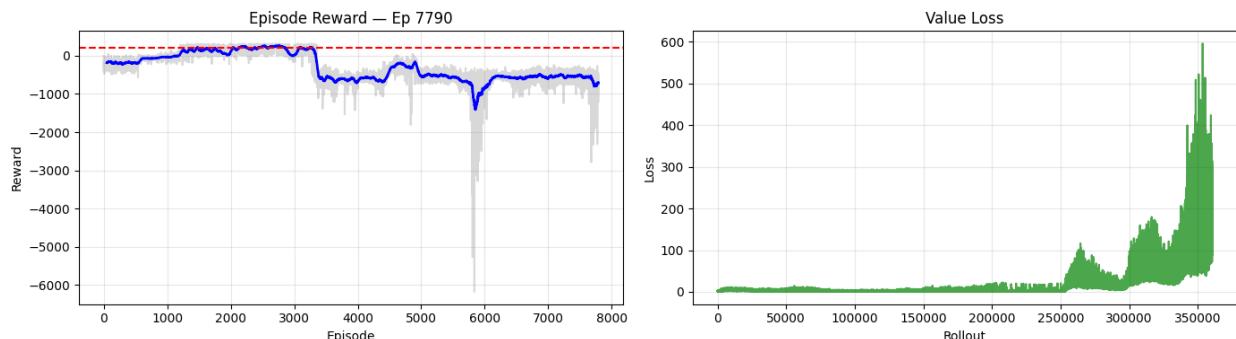
Final model: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_02_53
_32/dqn_ppo_dqn_42

Best model: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_02_53
_32/best_model

Checkpoints: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_02_53
_32/checkpoints

Best model stats: Reward: 283.74 +/- 17.39 | Success: 100% | Score (mean-std): 26
6.35 | @ 1,425,000 steps

```
=====
DQN | Seed 123
=====
```



```
'[Checkpoint] Episode 7700 saved'
```

```
'Eval @ 1500000 steps | Reward: -926.29 +/- 368.67 | Success: 0% | Score (mean-st
d): -1294.96 | Best: 259.82'
```

```
Output()
```

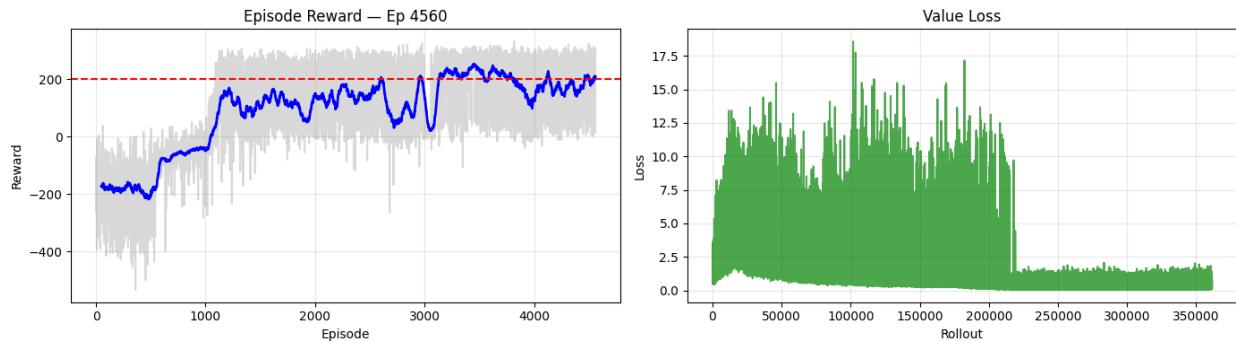
```
'Episode 4560 | Last 50 Ep - Mean: 209.7 | Std: 112.3 | Min: 0.6 | Max: 313.2 | S
uccess: 70%'
```

```

Training time: 47.0 min (2818 s)
Final model: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_03_41
_09/dqn_ppo_dqn_123
Best model: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_03_41
_09/best_model
Checkpoints: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_03_41
_09/checkpoints
Best model stats: Reward: 275.61 +/- 15.80 | Success: 100% | Score (mean-std): 25
9.82 | @ 875,000 steps

```

```
=====
DQN | Seed 3407
=====
```



```
'[Checkpoint] Episode 4500 saved'
```

```
'Eval @ 1500000 steps | Reward: 250.47 +/- 82.52 [SOLVED] | Success: 85% | Score
(mean-std): 167.95 | Best: 243.87'
```

```
Output()
```

```

Training time: 47.3 min (2839 s)
Final model: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_04_28
_07/dqn_ppo_dqn_3407
Best model: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_04_28
_07/best_model
Checkpoints: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_04_28
_07/checkpoints
Best model stats: Reward: 261.53 +/- 17.66 | Success: 100% | Score (mean-std): 24

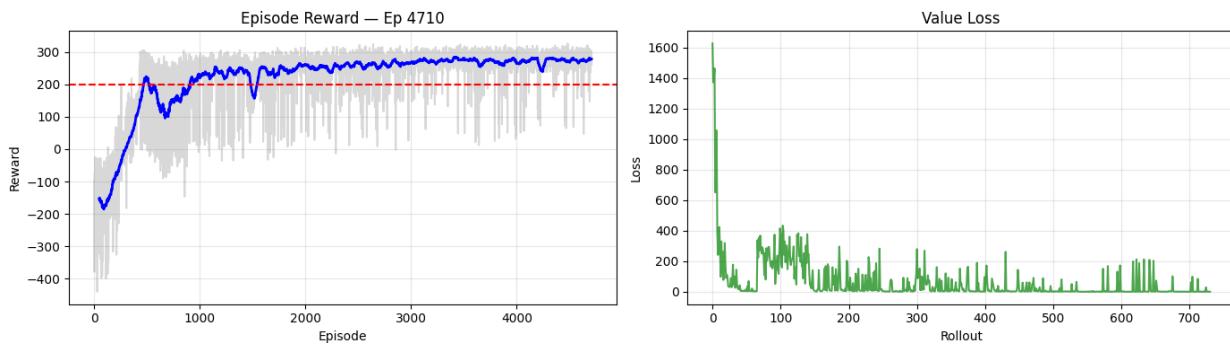
```

3.87 | @ 1,300,000 steps

DQN: All 3 seeds trained.

=====
PPO | Seed 42
=====

'Episode 4710 | Last 50 Ep – Mean: 278.3 | Std: 31.3 | Min: 146.7 | Max: 320.5 | Success: 96%'



'[Checkpoint] Episode 4700 saved'

'Eval @ 1500000 steps | Reward: 289.50 +/- 21.48 [SOLVED] | Success: 100% | Score

(mean-std): 268.02 | Best: 266.36 >> New best model!'

Output()

Training time: 20.2 min (1211 s)

Final model: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_15_27/dqn_ppo_ppo_42

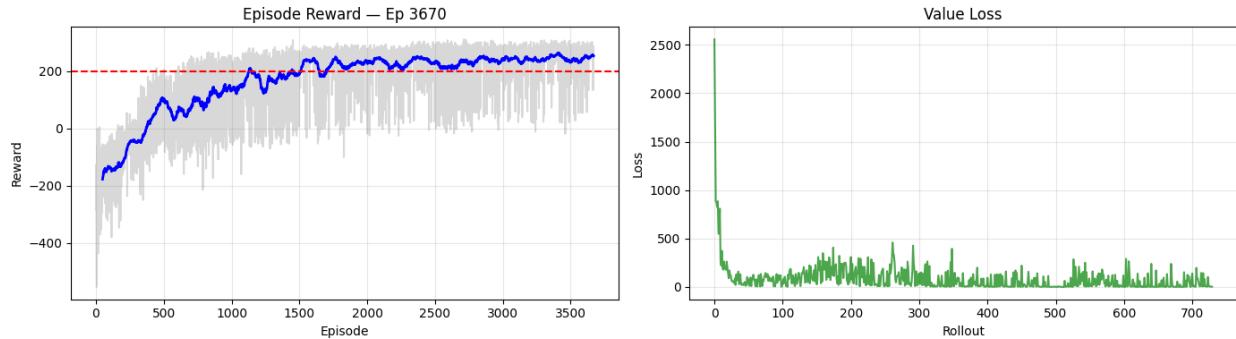
Best model: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_15_27/best_model

Checkpoints: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_15_27/checkpoints

Best model stats: Reward: 289.50 +/- 21.48 | Success: 100% | Score (mean-std): 268.02 | @ 1,500,000 steps

```
PPO | Seed 123
=====
```

```
'Episode 3670 | Last 50 Ep - Mean: 252.4 | Std: 46.9 | Min: 56.7 | Max: 301.1 | Success: 90%'
```



```
'[Checkpoint] Episode 3600 saved'
```

```
'Eval @ 1500000 steps | Reward: 257.63 +/- 35.30 [SOLVED] | Success: 95% | Score (mean-std): 222.33 | Best: 249.47'
Output()
```

```
Training time: 23.0 min (1380 s)
```

```
Final model: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_35_37/dqn_ppo_ppo_123
```

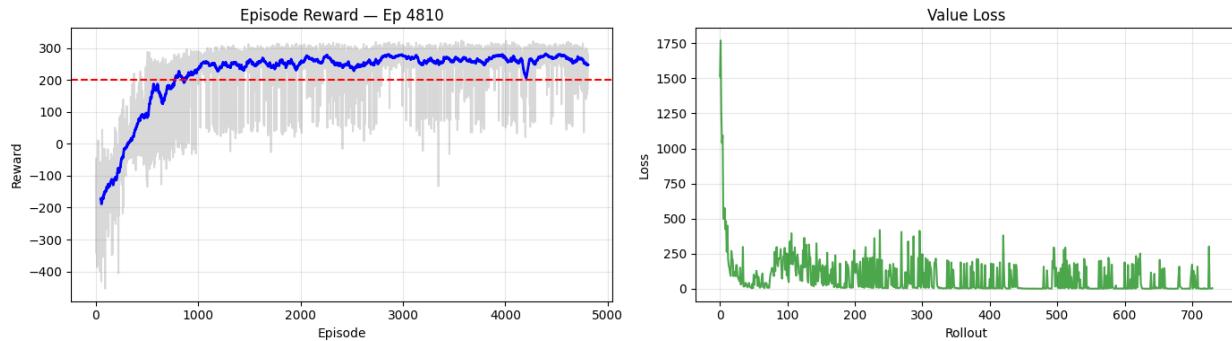
```
Best model: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_35_37/best_model
```

```
Checkpoints: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_35_37/checkpoints
```

```
Best model stats: Reward: 267.61 +/- 18.14 | Success: 100% | Score (mean-std): 249.47 | @ 1,300,000 steps
```

```
=====
PPO | Seed 3407
=====
```

'Episode 4810 | Last 50 Ep – Mean: 247.0 | Std: 54.3 | Min: 36.0 | Max: 299.4 | Success: 86%'



'[Checkpoint] Episode 4800 saved'

'Eval @ 1500000 steps | Reward: 265.17 +/- 52.61 [SOLVED] | Success: 95% | Score (mean-std): 212.56 | Best: 263.95'

Training time: 20.8 min (1248 s)

Final model: /home/legus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_58_38/dqn_ppo_ppo_3407

Best model: /home/legus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_58_38/best_model

Checkpoints: /home/legus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_58_38/checkpoints

Best model stats: Reward: 280.95 +/- 17.01 | Success: 100% | Score (mean-std): 263.95 | @ 1,375,000 steps

PPO: All 3 seeds trained.

=====

BEST MODEL SUMMARY (all algorithms x seeds)

=====

Algorithm	Seed	Mean Reward	Std Reward	Success	Score (mean-std)	@ Timestep
DQN	42	283.74	17.39	100%	266.35	1,425,000
DQN	123	275.61	15.80	100%	259.82	875,000
DQN	3407	261.53	17.66	100%	243.87	1,300,000
PPO	42	289.50	21.48	100%	268.02	1,500,000
PPO	123	267.61	18.14	100%	249.47	1,300,000
PPO	3407	280.95	17.01	100%	263.95	1,375,000

All training complete.

```
Saved DQN seed 42: 5424 episodes -> /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_02_53_32/training_log.npz  
Saved DQN seed 123: 7796 episodes -> /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_03_41_09/training_log.npz  
Saved DQN seed 3407: 4560 episodes -> /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_04_28_07/training_log.npz  
Saved PPO seed 42: 4710 episodes -> /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_15_27/training_log.npz  
Saved PPO seed 123: 3675 episodes -> /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_35_37/training_log.npz  
Saved PPO seed 3407: 4813 episodes -> /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_58_38/training_log.npz
```

***** TRAINING TIME SUMMARY *****

Timesteps per seed: 1,500,000 | Device: cpu

Algorithm Seed Time (s) Time (min)

DQN	42	2856	47.6
DQN	123	2818	47.0
DQN	3407	2839	47.3
PPO	42	1211	20.2
PPO	123	1380	23.0
PPO	3407	1248	20.8
DQN Mean		2838	47.3
PPO Mean		1280	21.3

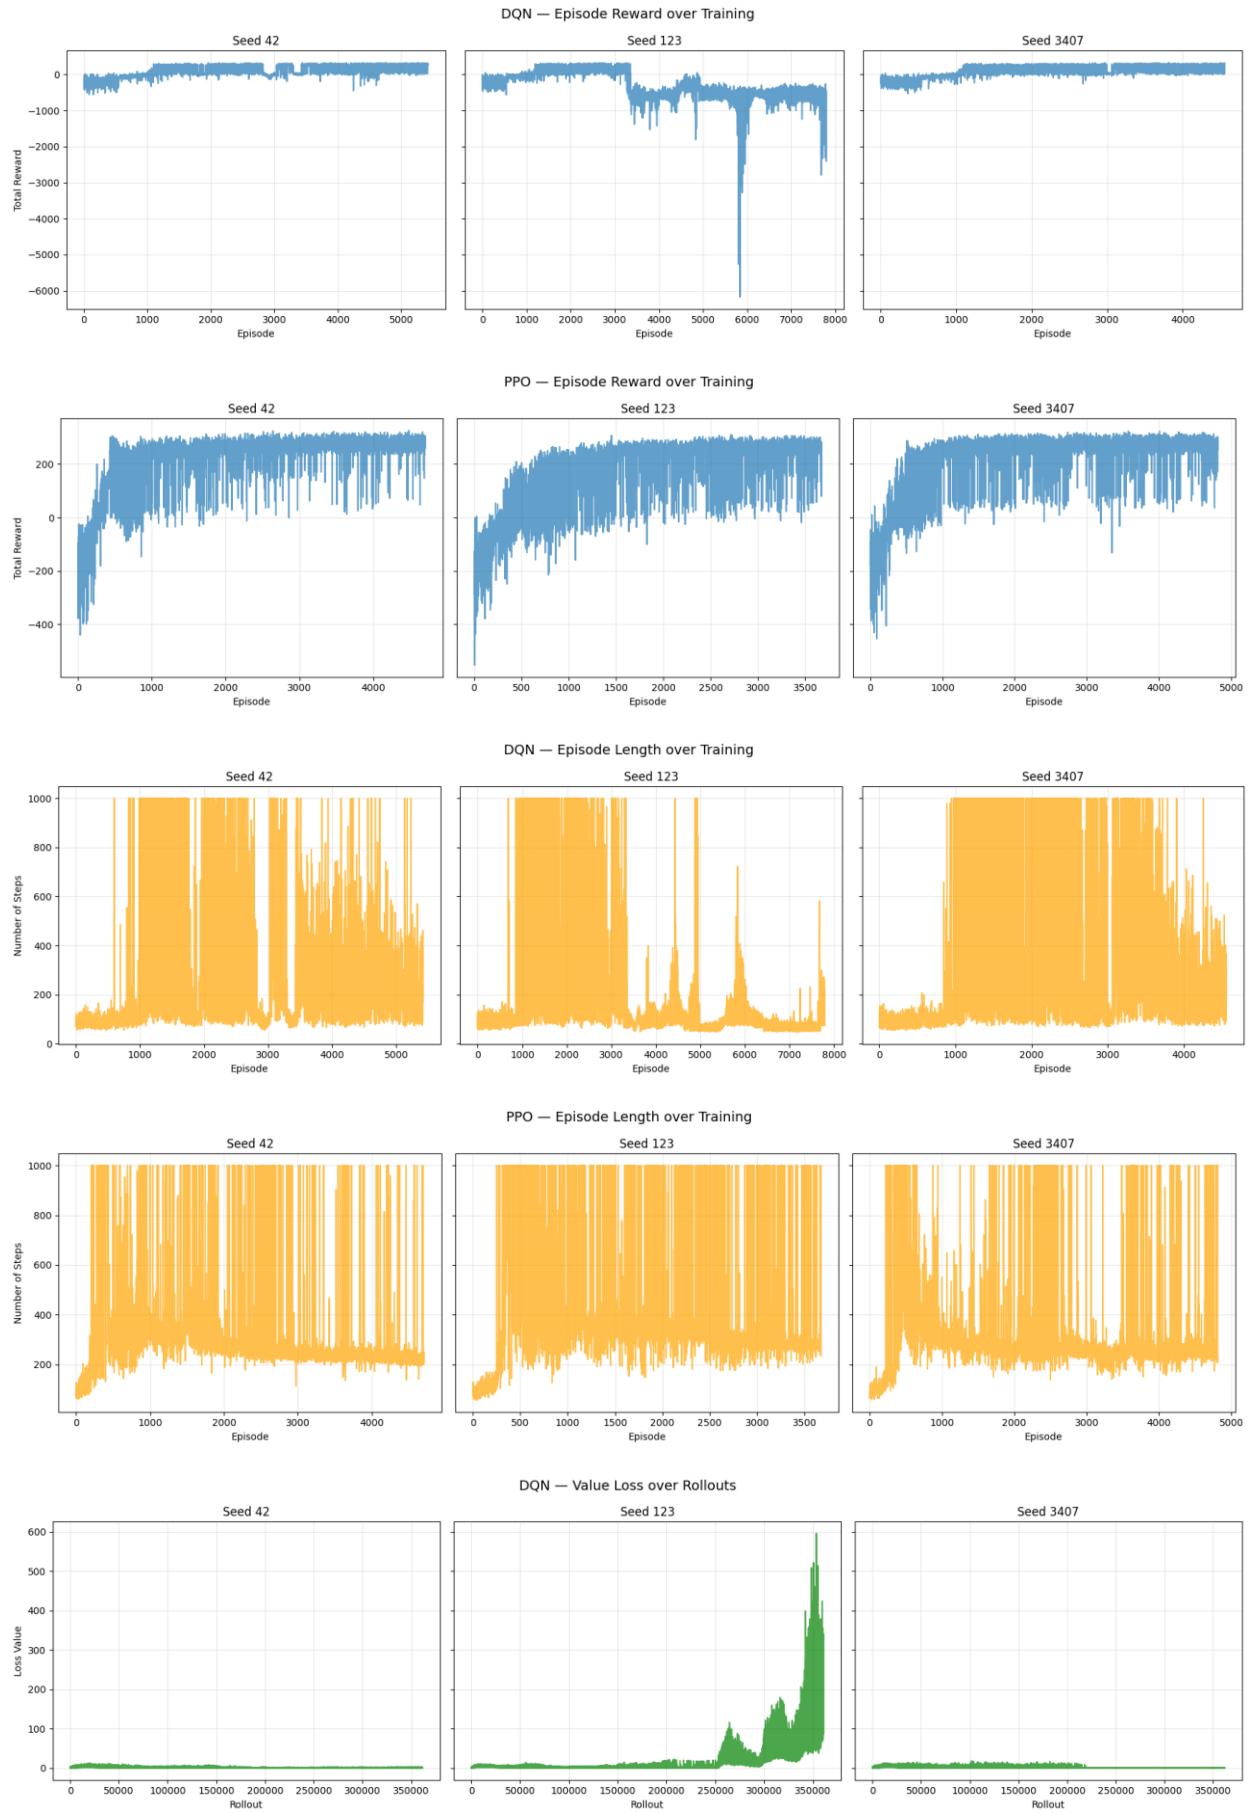
***** GRADIENT UPDATES SUMMARY *****

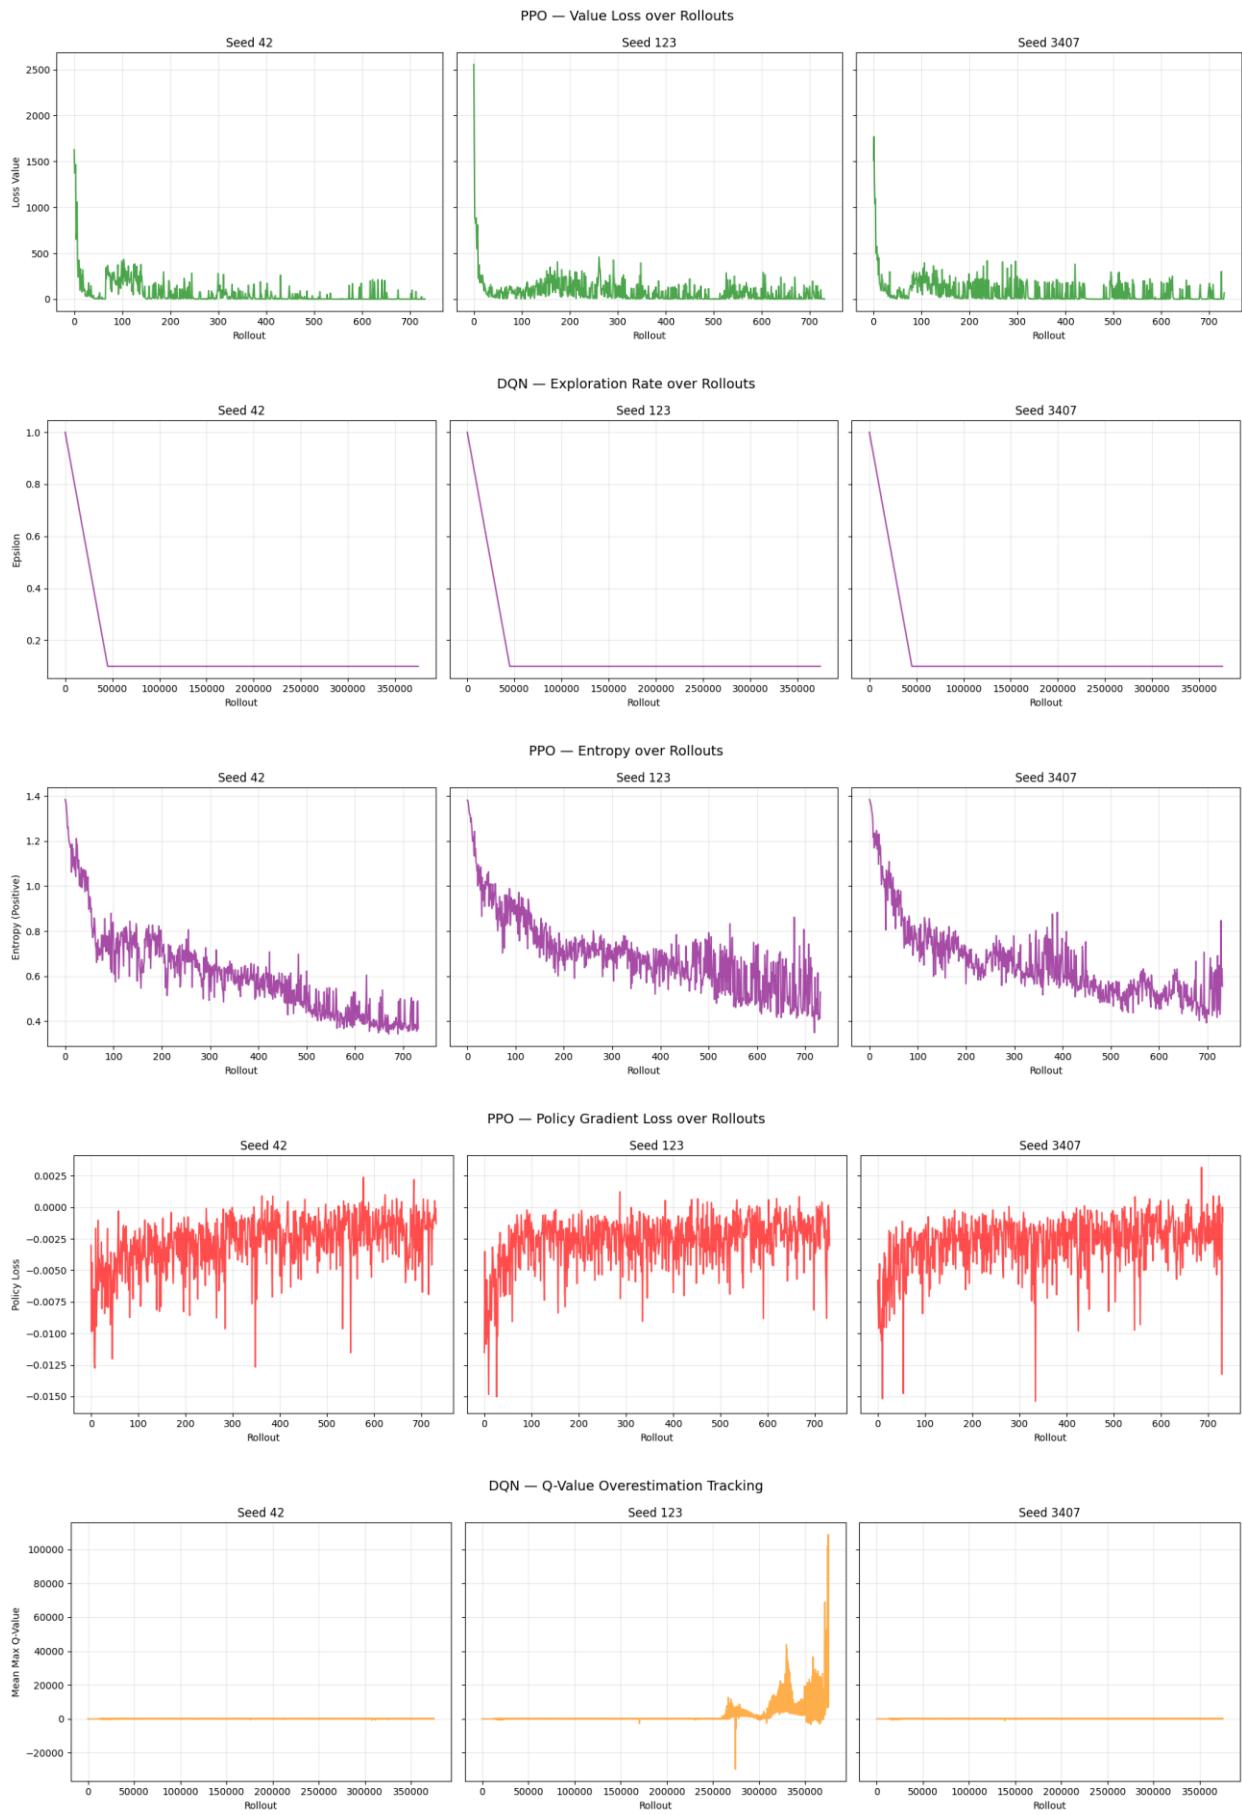
Algorithm Actual (from training) Analytical (computed) Total Env Steps Ratio (updates/steps)

DQN	1,449,996	1,450,000	1,500,000
0.967			
PPO	7,320	234,240	1,500,000
0.005			

DQN: train_freq=4, gradient_steps=4 -> 4 gradient updates per training call, called every 4 env steps

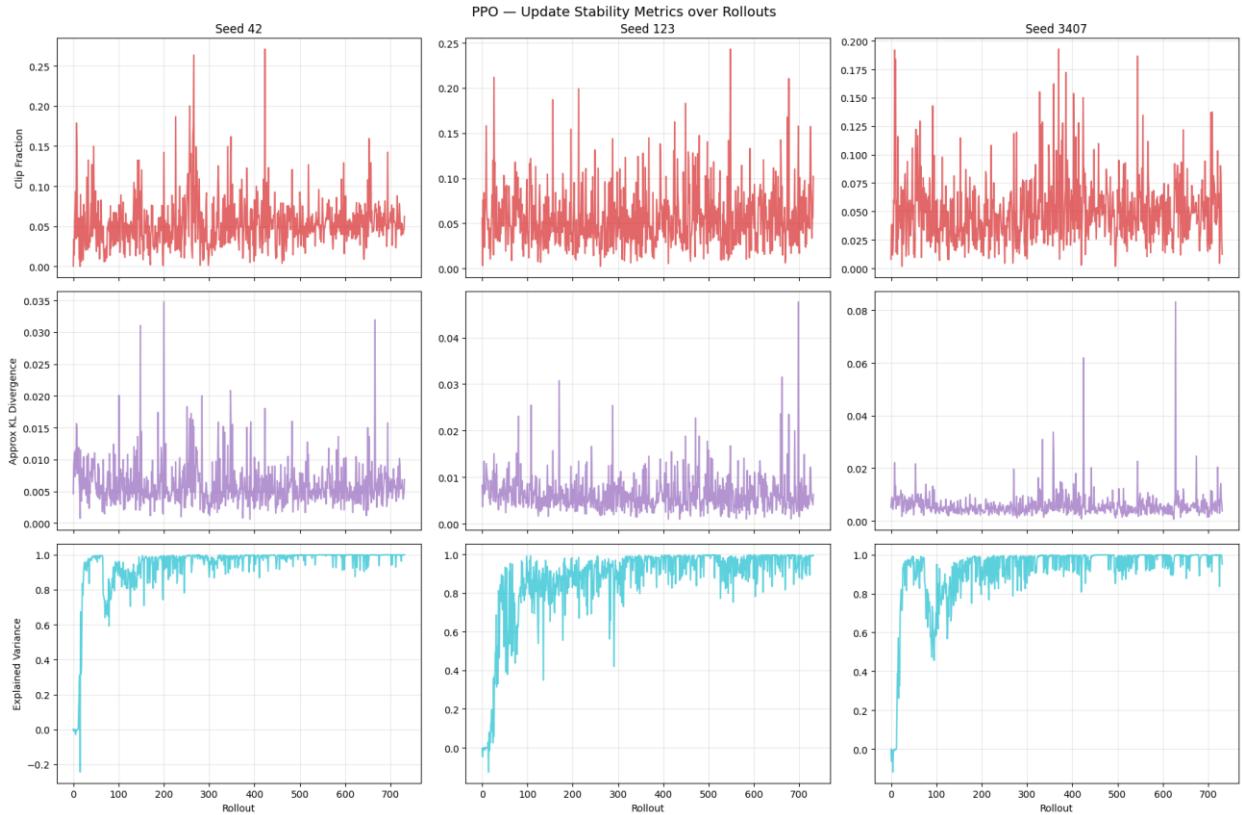
PPO: n_steps=2048, n_epochs=10, batch_size=64 -> 320 updates per rollout





Note: Steadily rising Q-values that diverge from actual returns indicate overestimation.

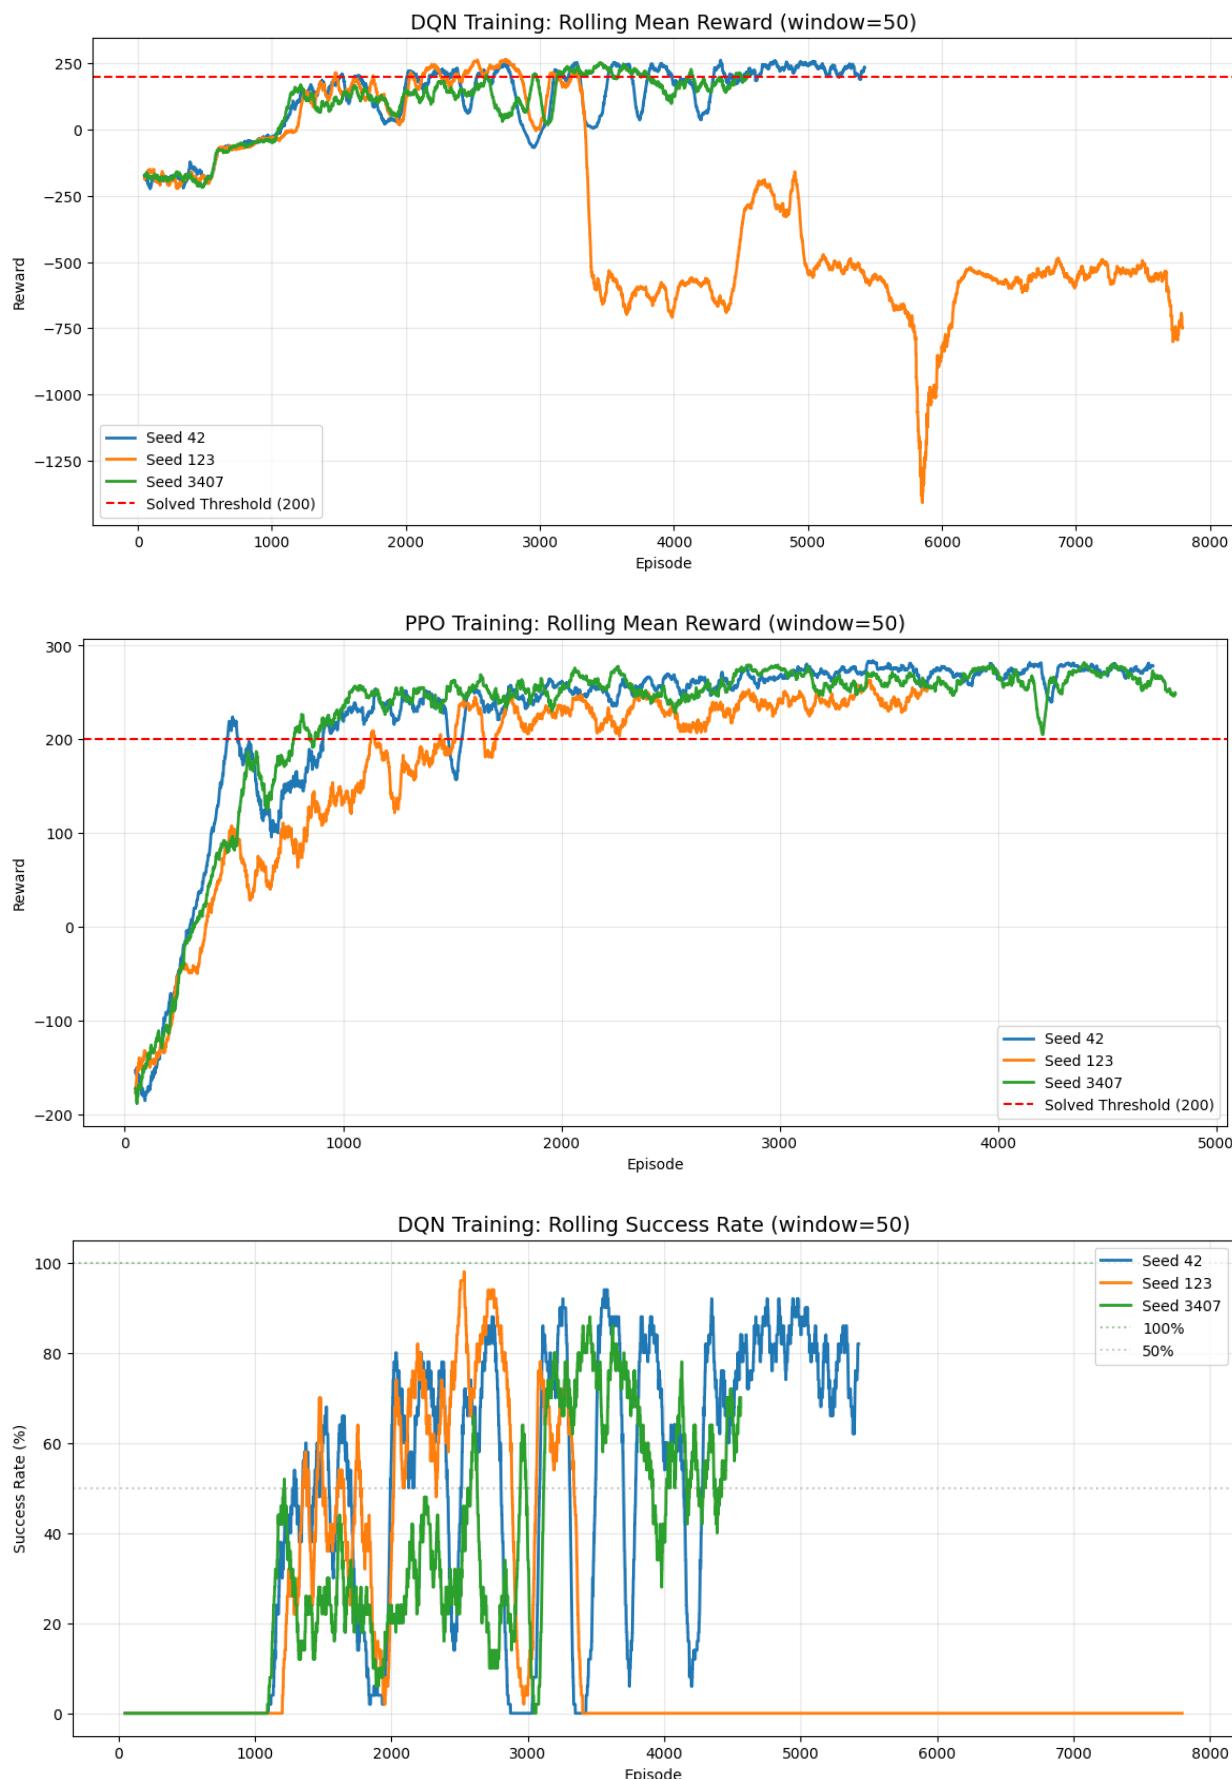
A stable or slowly growing curve suggests the target network is controlling overestimation.

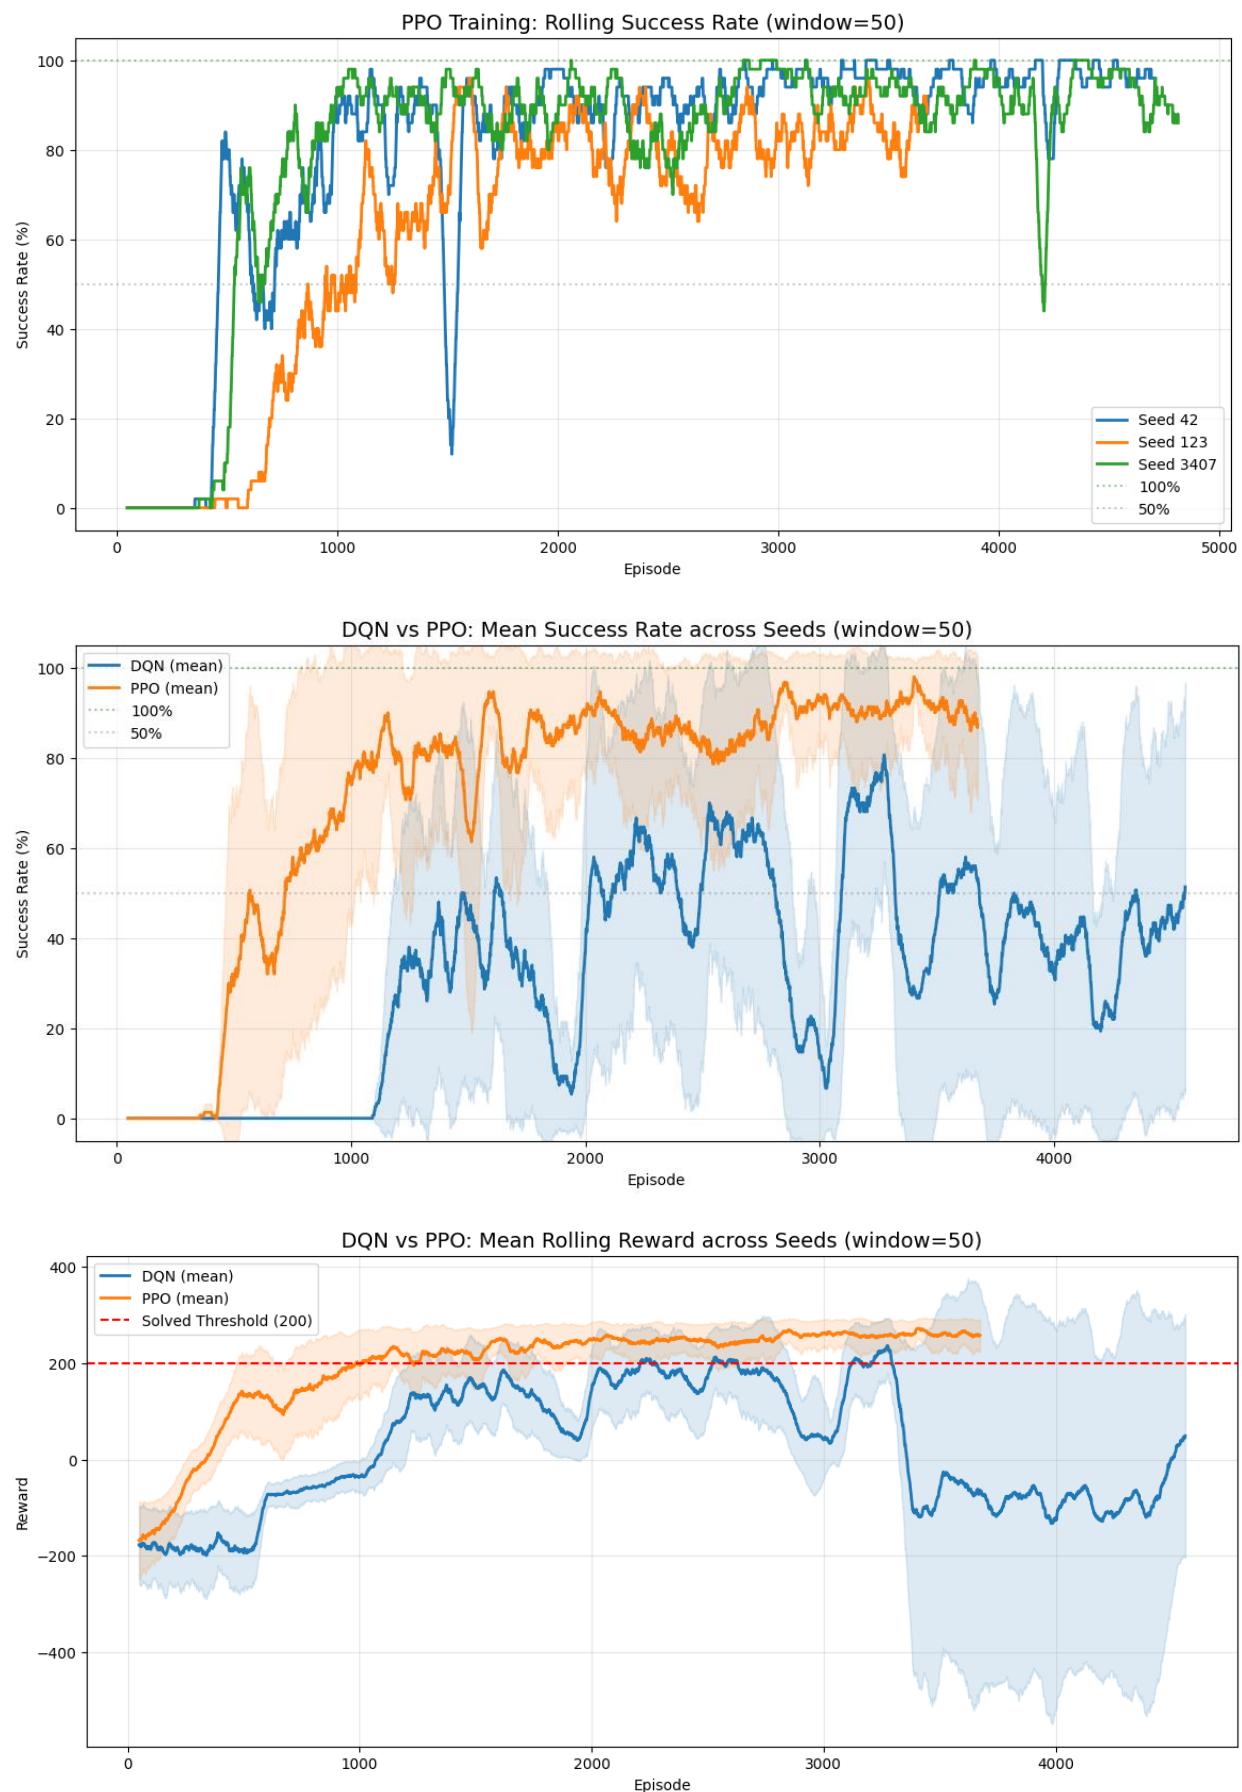


Clip Fraction: Fraction of policy updates clipped by PPO. High values suggest the policy is changing too fast.

Approx KL Divergence: KL divergence between old and new policy. Spikes indicate large policy shifts.

Explained Variance: How well the value function predicts returns. 1.0 = perfect, 0 = no better than mean.





4. Evaluation

The following tables summarize the evaluation performance of each algorithm across all seeds. Each model was evaluated over 20 deterministic episodes. The overall row aggregates all episodes across seeds.

```
Session: dqn_ppo
Algorithms: ['dqn', 'ppo']
Seeds: [42, 123, 3407]

Discovered best models:
  DQN seed 42: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_02_5
  3_32/best_model.zip
  DQN seed 123: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_03_
  41_09/best_model.zip
  DQN seed 3407: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_04
  _28_07/best_model.zip
  PPO seed 42: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_1
  5_27/best_model.zip
  PPO seed 123: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_
  35_37/best_model.zip
  PPO seed 3407: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05
  _58_38/best_model.zip

Discovered eval logs:
  DQN seed 42: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_02_5
  3_32/eval_log/evaluations.npz
  DQN seed 123: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_03_
  41_09/eval_log/evaluations.npz
  DQN seed 3407: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_04
  _28_07/eval_log/evaluations.npz
  PPO seed 42: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_1
  5_27/eval_log/evaluations.npz
  PPO seed 123: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_
  35_37/eval_log/evaluations.npz
  PPO seed 3407: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05
  _58_38/eval_log/evaluations.npz

Discovered training logs:
  DQN seed 42: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_02_5
  3_32/training_log.npz
  DQN seed 123: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_03_
  41_09/training_log.npz
  DQN seed 3407: /home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-23_04
  _28_07/training_log.npz
```

```

PPO seed 42: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_1
5_27/training_log.npz
PPO seed 123: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_
35_37/training_log.npz
PPO seed 3407: /home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-23_05_
58_38/training_log.npz

Loading and evaluating DQN seed 42 (best model)...
Loading and evaluating DQN seed 123 (best model)...
Loading and evaluating DQN seed 3407 (best model)...
DQN: evaluation complete.

Loading and evaluating PPO seed 42 (best model)...
Loading and evaluating PPO seed 123 (best model)...
Loading and evaluating PPO seed 3407 (best model)...
PPO: evaluation complete.

All evaluations complete.

```

4.1 Per-Algorithm Results

*** DQN MULTI-SEED EVALUATION SUMMARY ***

Episodes per seed: 20 | Total: 60

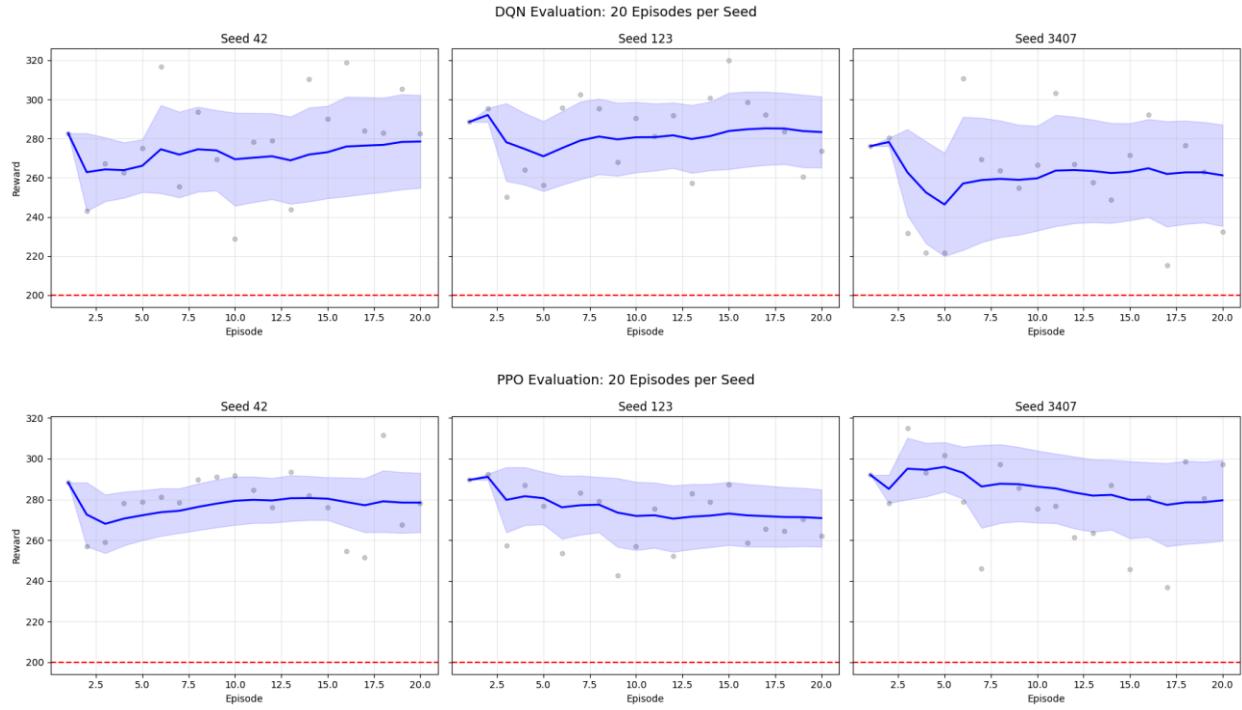
Seed	Mean Reward	Std Dev	Min Reward	Max Reward	Success Rate
42	278.51	23.70	228.67	318.97	100.0%
123	283.34	18.19	250.33	319.93	100.0%
3407	261.22	25.88	215.37	310.74	100.0%
Overall	274.36	24.72	215.37	319.93	100.0%

*** PPO MULTI-SEED EVALUATION SUMMARY ***

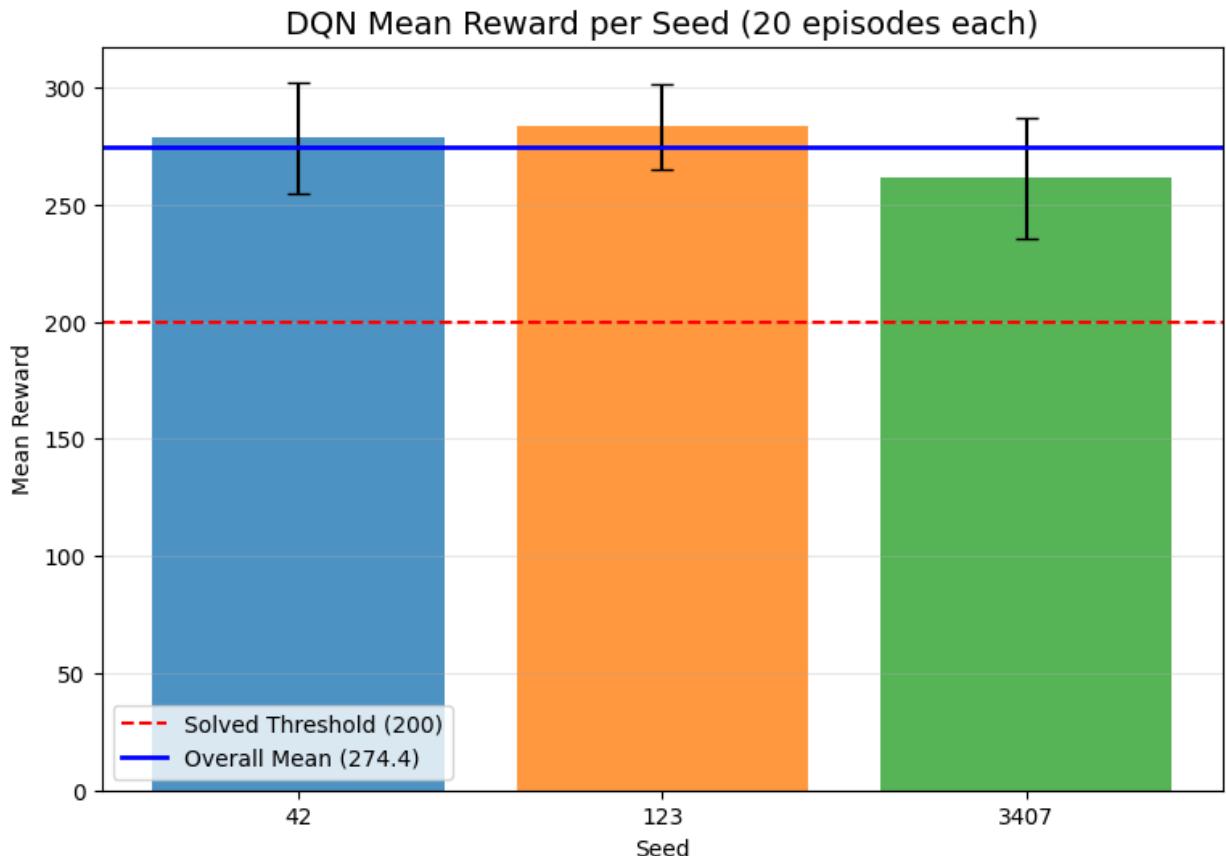
Episodes per seed: 20 | Total: 60

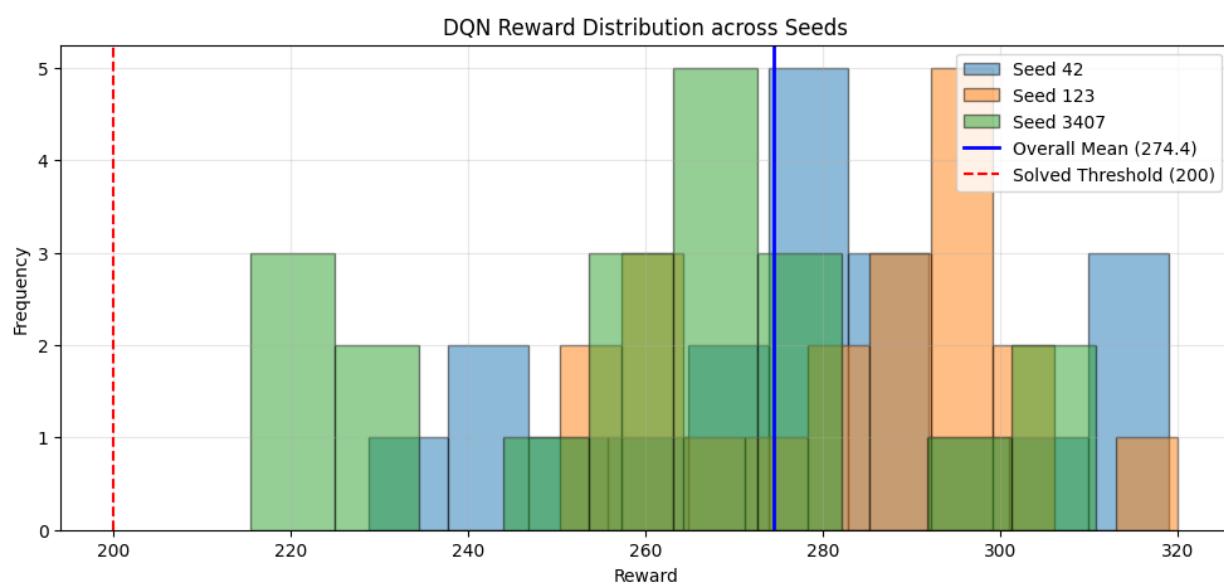
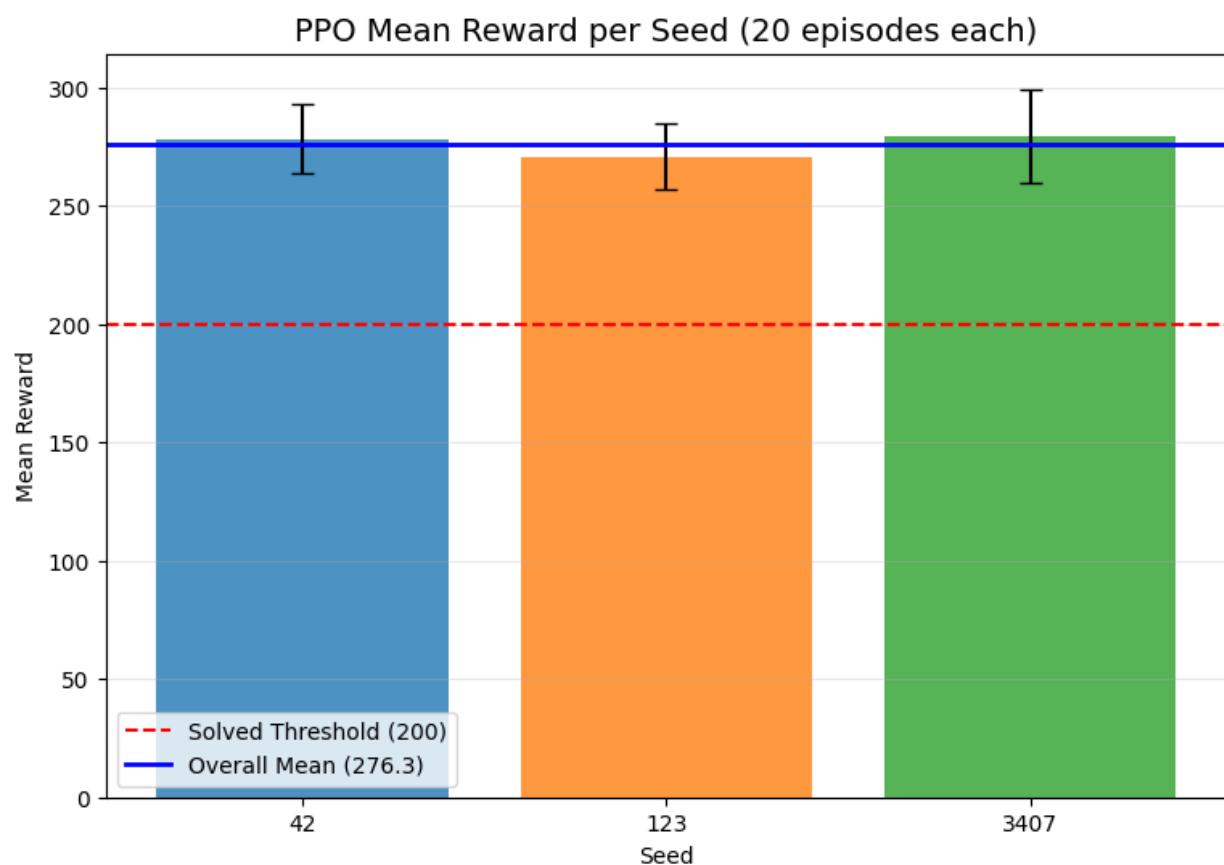
Seed	Mean Reward	Std Dev	Min Reward	Max Reward	Success Rate
42	278.43	14.56	251.45	311.70	100.0%
123	270.83	14.06	242.53	292.55	100.0%
3407	279.56	19.83	236.91	314.95	100.0%
Overall	276.27	16.81	236.91	314.95	100.0%

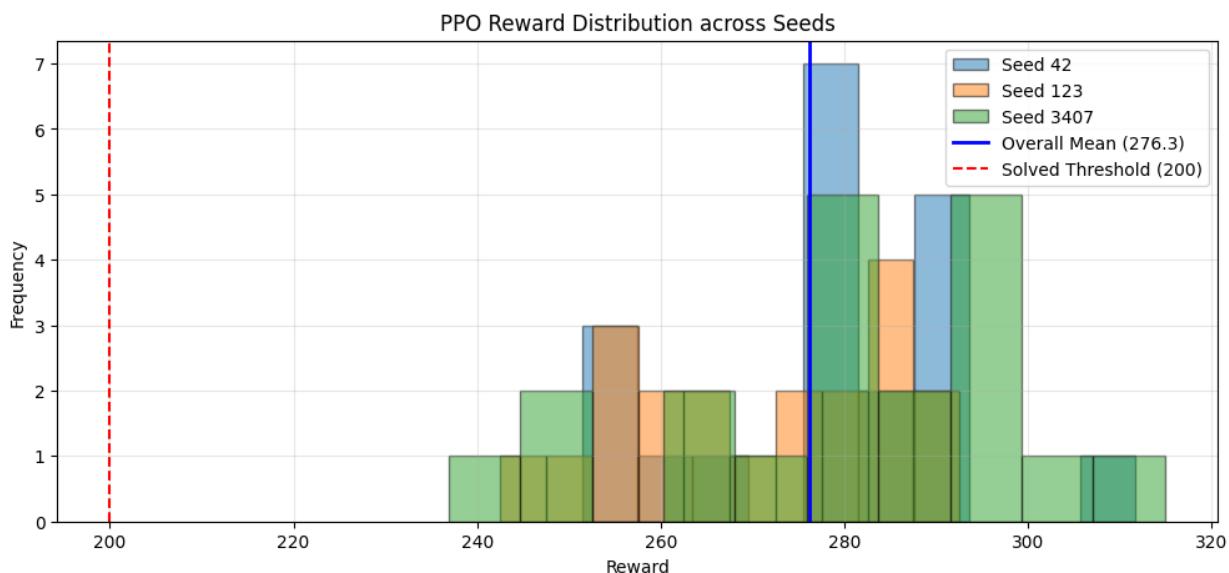
These convergence plots show the running mean reward across the 20 evaluation episodes for each seed. The shaded region represents one standard deviation. A flat, high running mean indicates consistent performance.



These convergence plots show the running mean reward across the 20 evaluation episodes for each seed. The shaded region represents one standard deviation. A flat, high running mean indicates consistent performance.







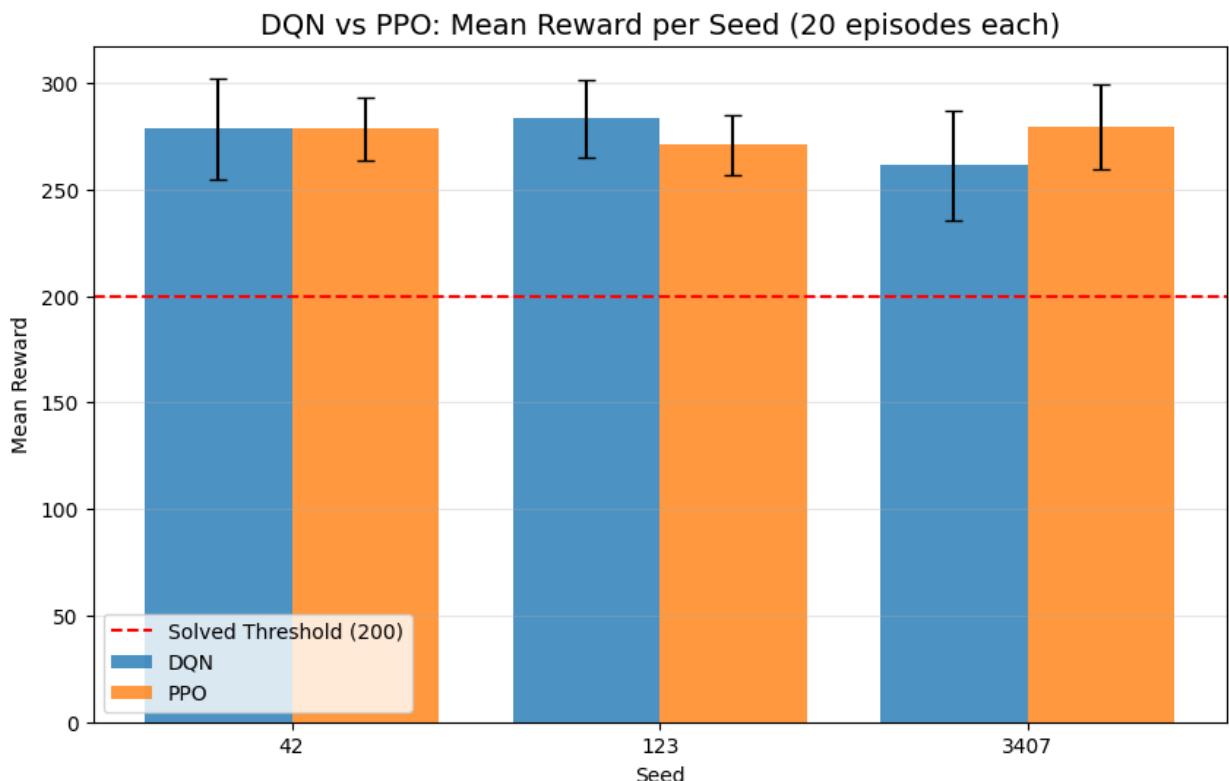
4.2 Cross-Algorithm Comparison

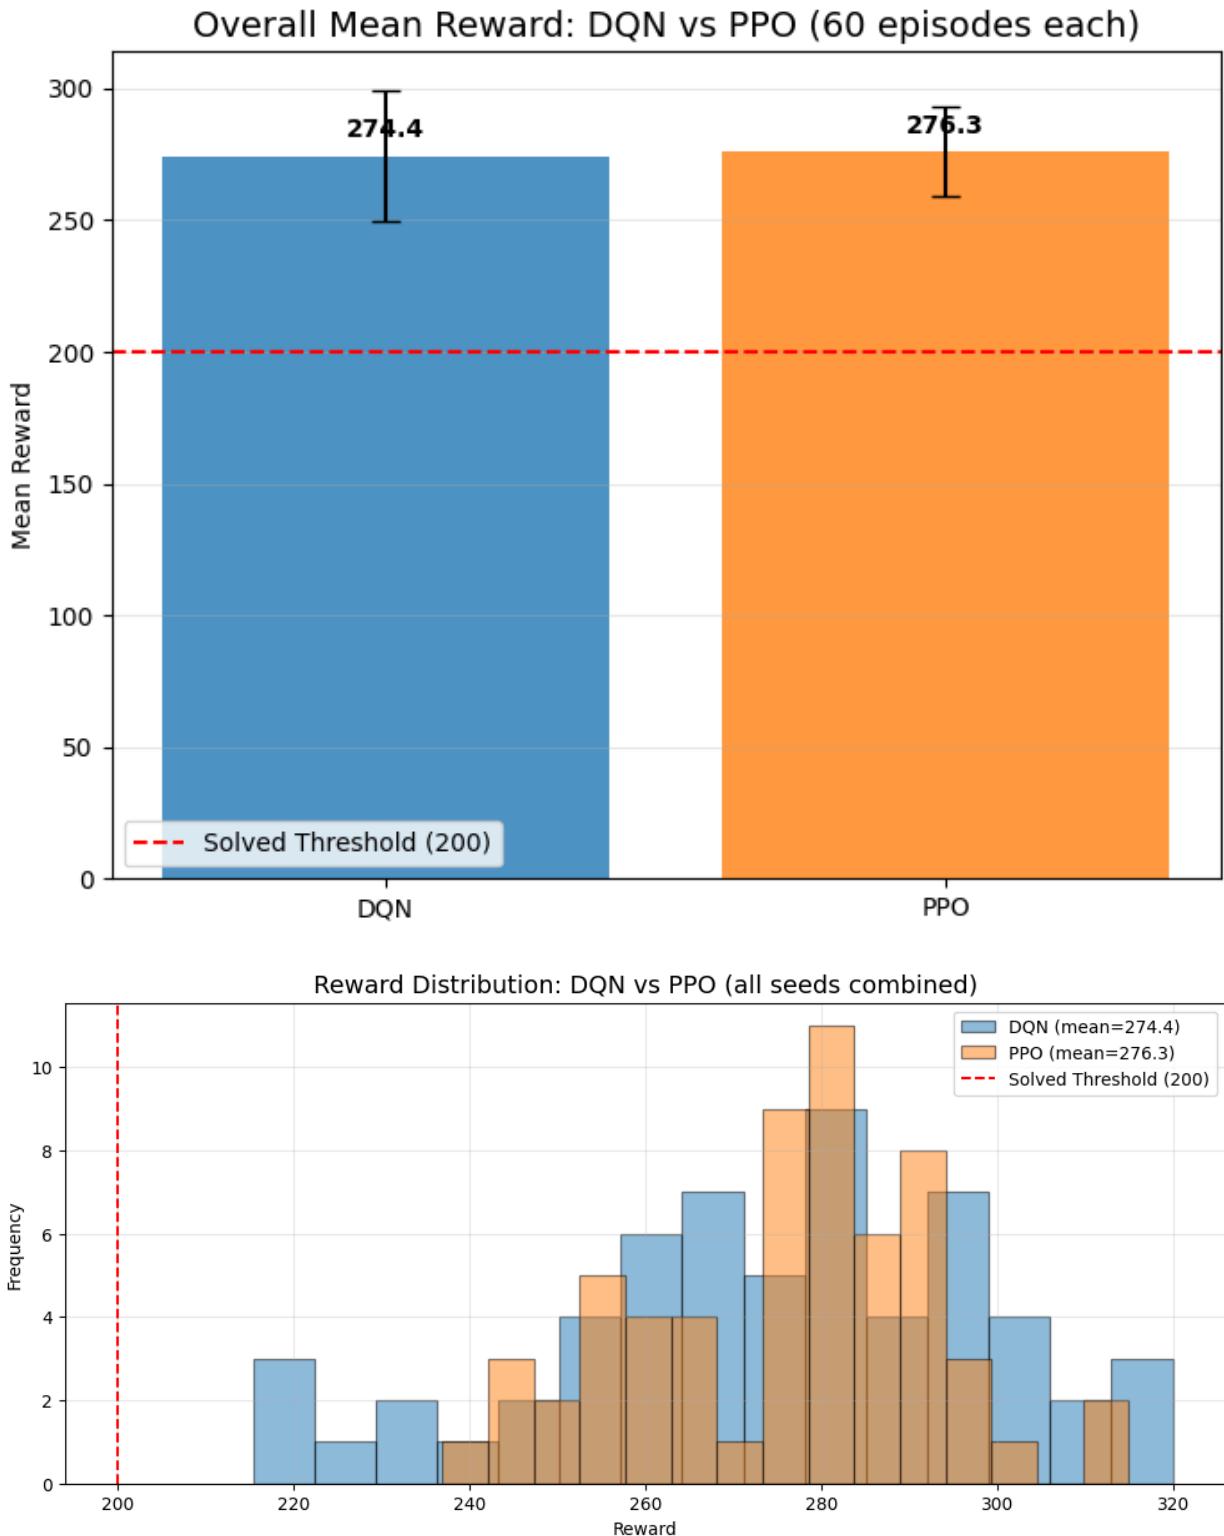
*** CROSS-ALGORITHM EVALUATION SUMMARY ***

Seeds: [42, 123, 3407] | Episodes per seed: 20

Total episodes per algorithm: 60

Algorithm	Mean Reward	Std Dev	Min Reward	Max Reward	Success Rate
DQN	274.36	24.72	215.37	319.93	100.0%
PPO	276.27	16.81	236.91	314.95	100.0%



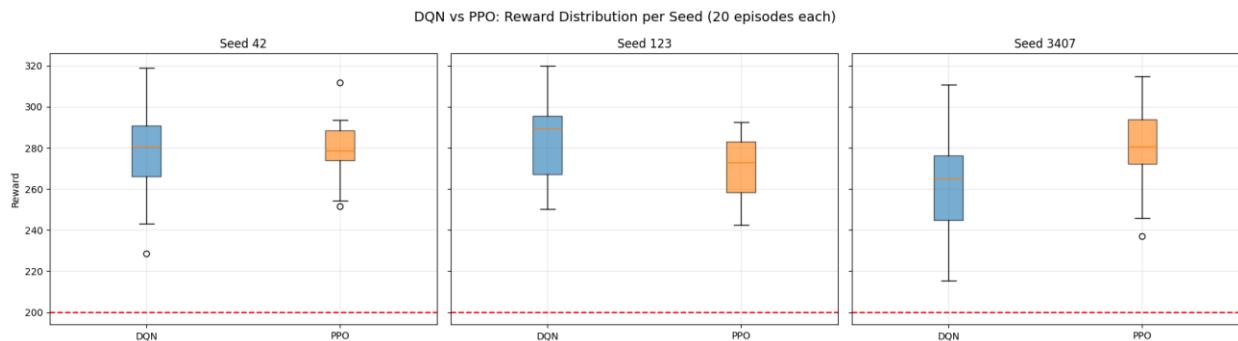


```
/tmp/ipykernel_6515/3208474440.py:9: MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.
```

```
    bp = ax.boxplot(data, labels=[a.upper() for a in algo_names], patch_artist=True)
```

```
/tmp/ipykernel_6515/3208474440.py:9: MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support
```

```
t for the old name will be dropped in 3.11.
bp = ax.boxplot(data, labels=[a.upper() for a in algo_names], patch_artist=True)
/tmp/ipykernel_6515/3208474440.py:9: MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.
bp = ax.boxplot(data, labels=[a.upper() for a in algo_names], patch_artist=True)
```



4.3 Statistical Significance

Two statistical tests assess whether the observed performance difference between DQN and PPO is significant. The Mann-Whitney U test compares reward distributions (non-parametric, no normality assumption). The Chi-squared test compares success rates as a proportion.

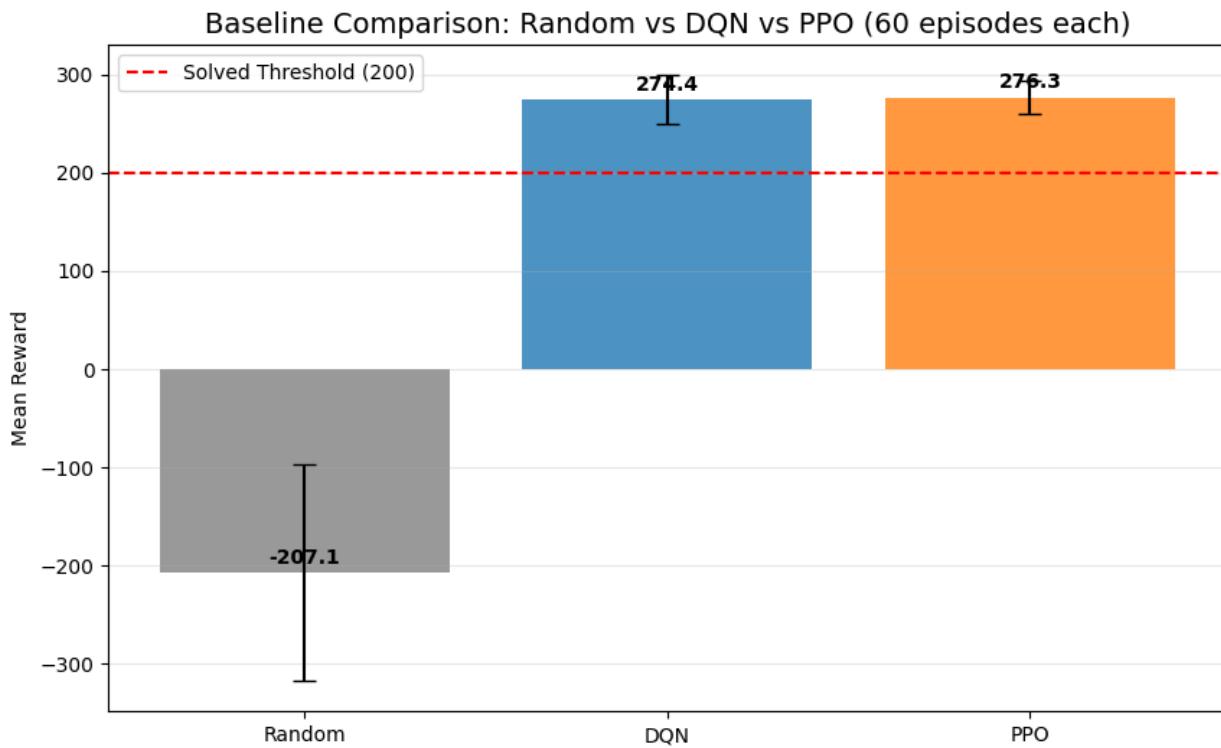
*** STATISTICAL SIGNIFICANCE TESTS ***					
Sample size per algorithm: 60 episodes (20 episodes x 3 seeds)					
Metric	DQN Value	PPO Value		Test Stat	
Statistic	p-value	Significant (p<0.05)			
Mean Reward	274.36	276.27		Mann-Whitney U	1
771.0	0.8811	No			
Success Rate (>=200)	100.0%	100.0%	Chi-squared (skipped: zero row/col)		
0.00	1.0000	No			

4.4 Baseline Comparison

```
Running random agent with seed 42...
Running random agent with seed 123...
Running random agent with seed 3407...
Random baseline evaluation complete.
```

```
*** BASELINE COMPARISON ***
```

Agent	Mean Reward	Std Dev	Min	Max	Success Rate
Random	-207.06	110.45	-416.10	15.45	0.0%
DQN	274.36	24.72	215.37	319.93	100.0%
PPO	276.27	16.81	236.91	314.95	100.0%

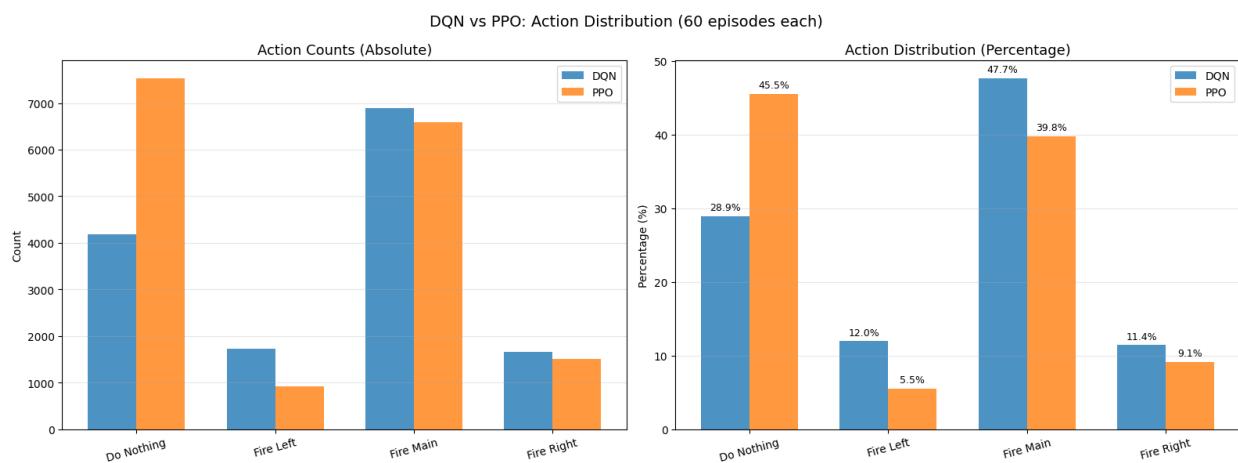


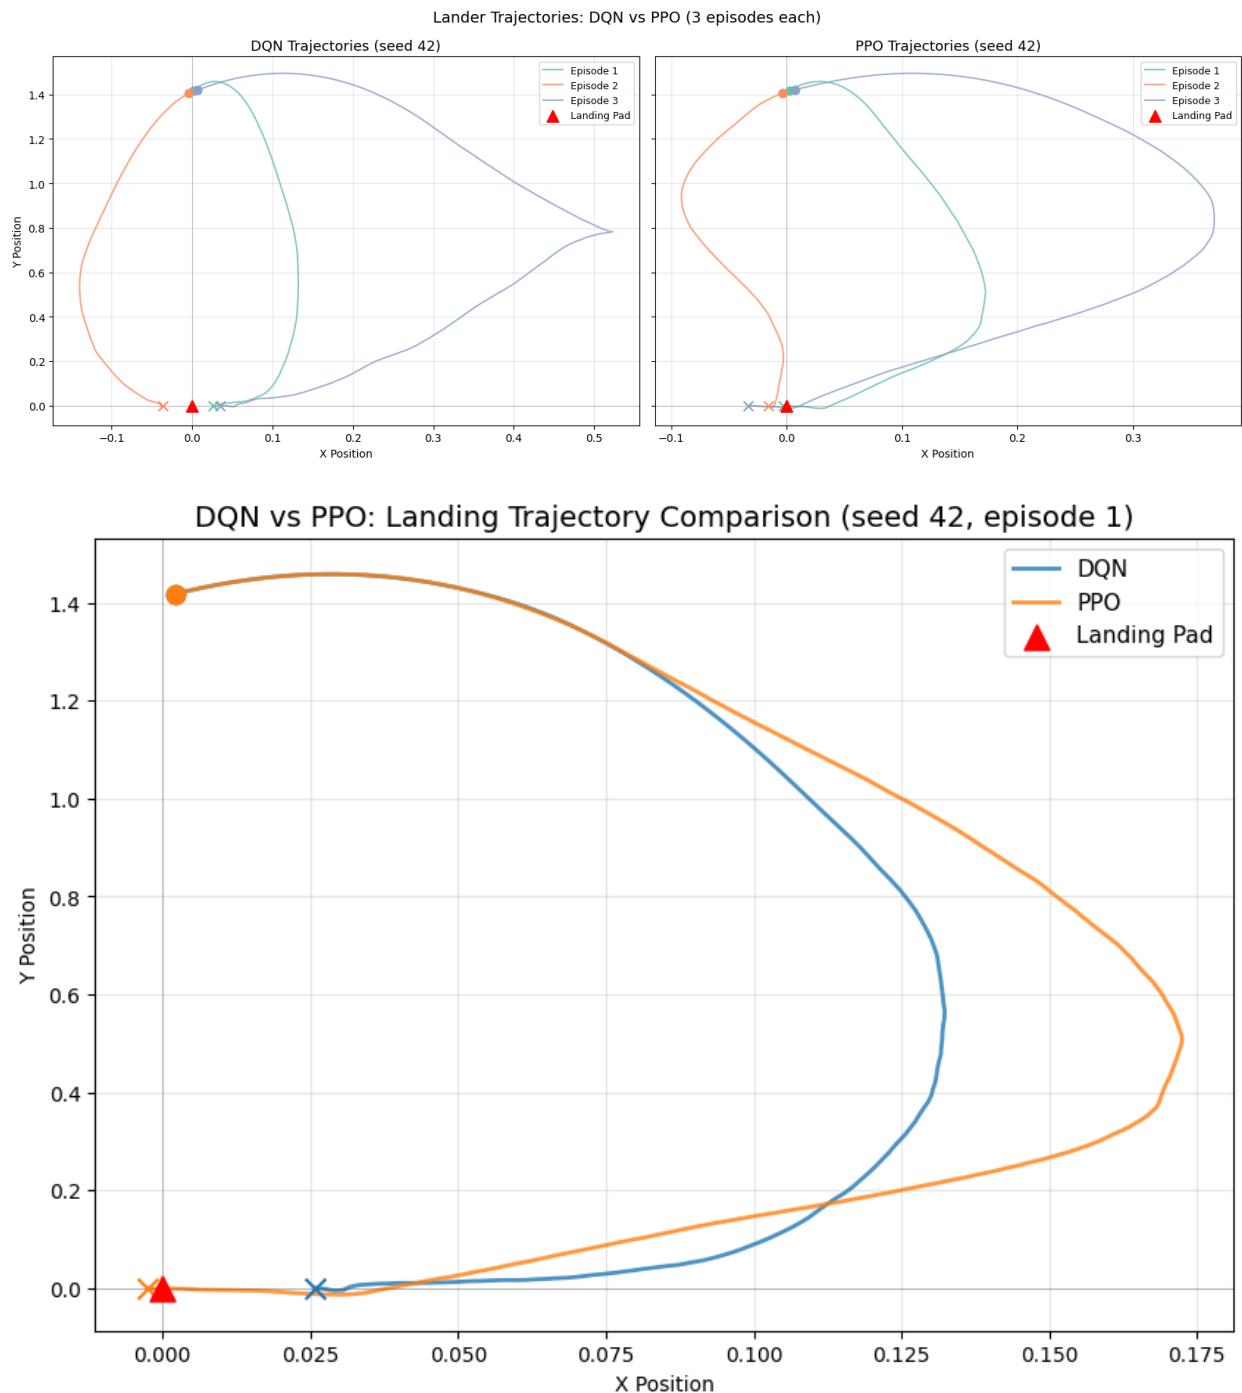
4.5 Agent Behavior Analysis

DQN: 14,451 total actions collected across 60 episodes

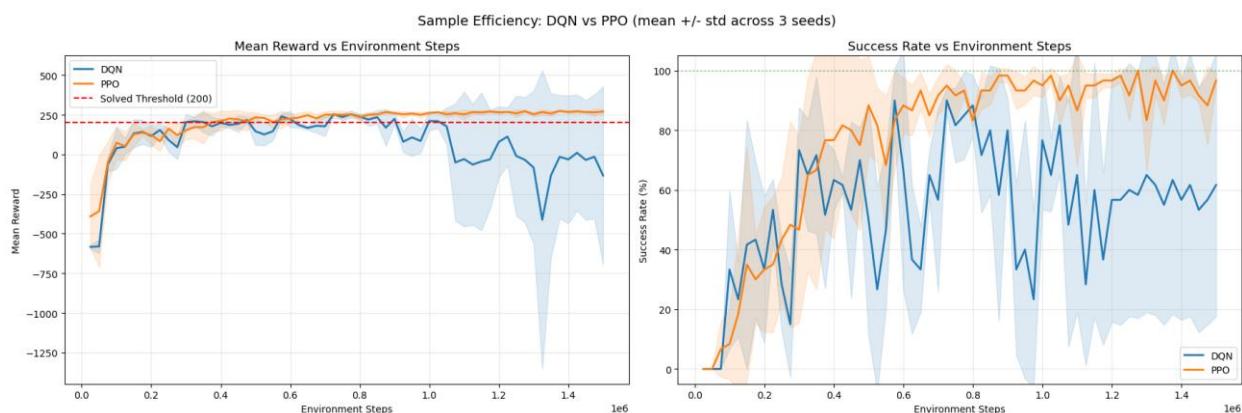
PPO: 16,544 total actions collected across 60 episodes

Behavior data collection complete.





4.6 Sample Efficiency Analysis



```
*** SAMPLE EFFICIENCY: FIRST SOLVED Timestep ***
Solved = mean evaluation reward >= 200
Evaluated every 25,000 timesteps with 20 episodes
```

Algorithm Seed First Solved at Step

DQN	42	300,000
DQN	123	300,000
DQN	3407	150,000
DQN	Mean	250,000
PPO	42	150,000
PPO	123	525,000
PPO	3407	325,000
PPO	Mean	333,333

4.7 Training Milestone Timeline

This chart overlays the training reward curves with vertical milestone markers showing when each algorithm reached key performance thresholds. The background curves show the rolling mean reward (window=100 episodes) plotted against cumulative environment steps, reconstructed from the per-episode training logs. Three milestones are marked per algorithm:

- Best model (solid line): the checkpoint selected by the combined metric (mean - std) during training
- First solved (dashed line): first point where the rolling 100-episode mean reward >= 200
- First all-pass (dotted line): first point where all 100 consecutive episodes scored >= 200

Discovered	training	logs:	DQN	seed	
/home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-					42:
23_02_53_32/training_log.npz			DQN	seed	123:
/home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-					
23_03_41_09/training_log.npz			DQN	seed	3407:
/home/logus/env/iscte/taap_p2/report/..../models/dqn/2026-02-					
23_04_28_07/training_log.npz			PPO	seed	42:
/home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-					
23_05_15_27/training_log.npz			PPO	seed	123:
/home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-					
23_05_35_37/training_log.npz			PPO	seed	3407:
/home/logus/env/iscte/taap_p2/report/..../models/ppo/2026-02-					
23_05_58_38/training_log.npz					

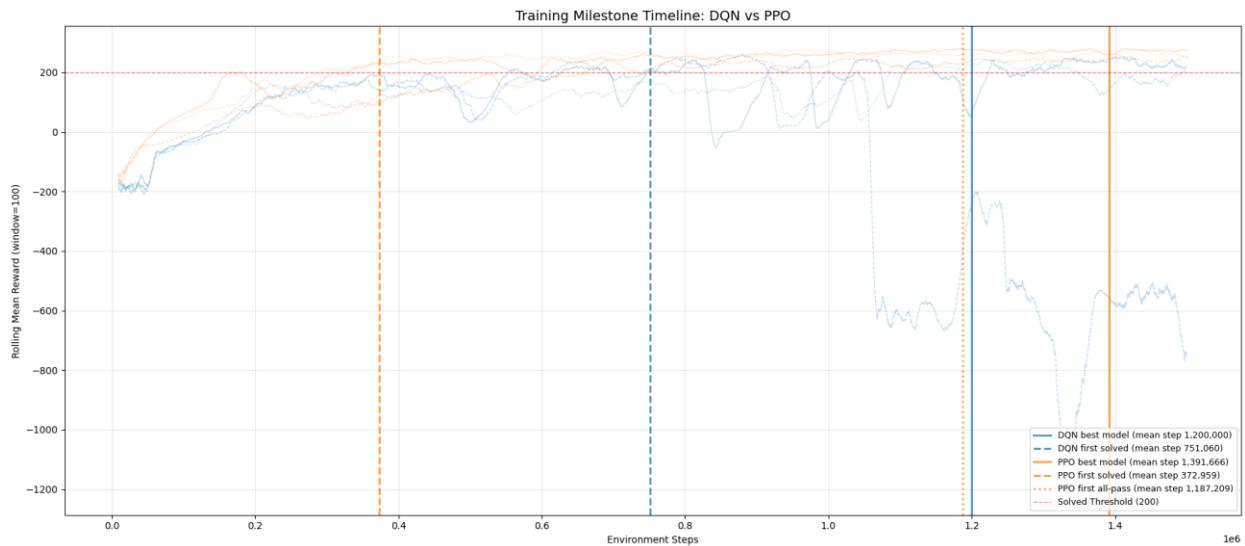


Figure 14 - Training milestone timeline showing rolling mean reward curves (per-seed, light lines) with vertical markers for best model selection, first solved (rolling 100-episode mean ≥ 200), and first all-pass (all 100 consecutive episodes ≥ 200) milestones.

*** TRAINING MILESTONES PER SEED ***				
Rolling window: 100 episodes				
Solved threshold: mean reward ≥ 200				
All-pass threshold: every episode in window ≥ 200				
Algorithm Seed First Solved (rolling mean ≥ 200) First All-Pass (all 100 ≥ 200)				
d	DQN	42	553,533	Not reached
d	DQN	123	614,569	Not reached
d	DQN	3407	1,085,079	Not reached
2	PPO	42	166,110	1,407,53
d	PPO	123	662,791	Not reached
6	PPO	3407	289,977	966,88

4.8 Final vs Best Model Comparison

```
Evaluating DQN seed 42 (final model)...
Evaluating DQN seed 123 (final model)...
Evaluating DQN seed 3407 (final model)...
Evaluating PPO seed 42 (final model)...
Evaluating PPO seed 123 (final model)...
```

```
Evaluating PPO seed 3407 (final model)...
```

```
Final model evaluation complete.
```

```
*** BEST MODEL vs FINAL MODEL ***
```

```
Episodes per evaluation: 20
```

Algorithm	Seed	Best	Mean	Best	Success	Final	Mean	Final	Success	Delta	Mean
DQN	42	278.51		100%		266.82		95%		+11.69	
DQN	123	283.34		100%		-935.49		0%		+1218.83	
DQN	3407	261.22		100%		286.92		100%		-25.70	
PPO	42	278.43		100%		276.39		95%		+2.04	
PPO	123	270.83		100%		247.72		85%		+23.11	
PPO	3407	279.56		100%		267.67		95%		+11.88	

```
Positive Delta Mean = best model selection improved over end-of-training snapshot.
```

4.9 Hyperparameter Exploration Summary

```
*** DQN Final Hyperparameters ***
```

Parameter	Value
policy	MlpPolicy
learning_rate	<function linear_schedule.<locals>.func at 0x75da6d15cc20>
learning_starts	50000
buffer_size	750000
batch_size	128
gamma	0.99
exploration_fraction	0.12
exploration_final_eps	0.1
target_update_interval	250
train_freq	4
gradient_steps	4
policy_kwargs	{'net_arch': [256, 256]}
device	cpu
total_timesteps	1,500,000
device	cpu
policy	MlpPolicy

```
*** PPO Final Hyperparameters ***
```

Parameter	Value
learning_rate	0.00025
n_steps	2048

```

batch_size          64
n_epochs           10
gamma              0.999
gae_lambda         0.95
ent_coef            0.01
clip_range          0.2
total_timesteps   1,500,000
device              cpu
policy              MlpPolicy

```

*** DQN HYPERPARAMETER EXPLORATION ***

Config	Key Changes	Success Rate	Mean Reward	Notes
A (baseline)	lr=6.3e-4, buffer=100k, eps_final=0.01, explore=0.5, soft updates (tau=0.005)	84.0%	245.47	Slow exploration decay, soft target updates
B (Zoo-style)	lr=linear(6.3e-4), buffer=1M, eps_final=0.1, explore=0.12, hard updates (250)	96.0%	263.02	Short exploration, high final epsilon, hard target updates
C (final)	lr=linear(6.3e-4), buffer=750k, eps_final=0.1, explore=0.12, hard updates (250), learning_starts=50k TBD (lab011) TBD (lab011) Added learning_starts buffer fill, gradient_steps=4 explicit			

*** PPO HYPERPARAMETER EXPLORATION ***

Config	Key Chan				
ges	Success	Rate	Mean	Reward	Notes
A (final)	lr=2.5e-4, gamma=0.999, n_steps=2048, batch=64, n_epochs=10, ent_coef=0.01	98.0%	269.21	Standard config, single environment	

Note: Update Config C results and add any additional configs tested during lab011 runs.

4.10 Visualizations

```

Generating GIF for DQN seed 42 (best model)...
Saved: /home/logus/env/iscte/taap_p2/report/outputs_dqn/dqn_seed42.gif

```

```
<IPython.core.display.Image object>

Generating GIF for DQN seed 123 (best model)...
Saved: /home/logus/env/iscte/taap_p2/report/outputs_dqn/dqn_seed123.gif

<IPython.core.display.Image object>

Generating GIF for DQN seed 3407 (best model)...
Saved: /home/logus/env/iscte/taap_p2/report/outputs_dqn/dqn_seed3407.gif

<IPython.core.display.Image object>

Generating GIF for PPO seed 42 (best model)...
Saved: /home/logus/env/iscte/taap_p2/report/outputs_ppo/ppo_seed42.gif

<IPython.core.display.Image object>

Generating GIF for PPO seed 123 (best model)...
Saved: /home/logus/env/iscte/taap_p2/report/outputs_ppo/ppo_seed123.gif

<IPython.core.display.Image object>

Generating GIF for PPO seed 3407 (best model)...
Saved: /home/logus/env/iscte/taap_p2/report/outputs_ppo/ppo_seed3407.gif
```

5. Experimental Setup

```
Environment: LunarLander-v3
Observation space: Box([-2.5      -2.5      -10.      -10.      -6.28318
55 -10.
-0.      -0.      ], [ 2.5      2.5      10.      10.      6.283185
5 10.
1.      1.      ], (8,), float32)
Action space: Discrete(4)
Wind enabled: False

Sample observation: [ 0.00229702  1.4181306   0.2326471   0.3204666  -0.00265488
-0.05269805
0.      0.      ]
Observation labels: [x, y, vx, vy, angle, angular_vel, left_leg, right_leg]

Python: 3.12.3
PyTorch: 2.10.0+cu130
Stable-Baselines3: 2.7.1
Gymnasium: 1.2.3
NumPy: 2.4.2
Device: cpu
CUDA: 13.0
```

6. Conceptual Discussion

This section presents a theoretical discussion of the core conceptual differences between value-based and policy-based reinforcement learning methods, the sources of instability in deep RL, and the exploration mechanisms used in DQN and PPO.

6.1 Value-Based vs Policy-Based Learning

Reinforcement Learning algorithms can be categorised according to what is parameterised and optimised.

Value-Based Learning

Value-based methods approximate a value function, typically the action-value function:

$$Q^\pi(s, a) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s, a_0 = a \right]$$

The optimal policy is derived implicitly:

$$\pi(s) = \operatorname{argmax}_a Q(s, a)$$

Deep Q-Networks (DQN) approximate $Q(s, a; \theta)$ and update parameters using the Bellman optimality equation:

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

This introduces **bootstrapping**, since the target depends on the model's own predictions.

Key characteristics:

- Off-policy learning
- Bootstrapping
- Policy extracted via maximisation

However, the use of the max operator combined with function approximation introduces instability.

Policy-Based Learning

Policy-based methods parameterise the policy directly:

$$\pi_\theta(a|s)$$

and optimise the expected return:

$$J(\theta) = \mathbb{E}_{\pi_\theta}[R]$$

Using the policy gradient theorem:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta}[\nabla_\theta \log \pi_\theta(a|s) A^\pi(s, a)]$$

These methods:

- Learn the policy directly
- Avoid explicit maximisation over value estimates
- Naturally handle continuous action spaces

Because they do not rely on a max operator over Q-values, they tend to be structurally more stable in deep settings.

6.2 Overestimation and Instability in Value-Based Algorithms

Overestimation Bias

Overestimation arises from the maximisation step $\max_a Q(s, a)$. If Q-value estimates contain zero-mean noise:

$$\mathbb{E}[\max(X_i)] \geq \max(\mathbb{E}[X_i])$$

Thus, even unbiased estimators lead to positively biased maximum estimates.

In DQN, the same network selects the action (argmax) and evaluates the action. This coupling amplifies overestimation bias.

Structural Instability: The Deadly Triad

Deep value-based RL combines:

1. Function approximation
2. Bootstrapping
3. Off-policy learning

This combination is known as the **deadly triad**, which may cause divergence or unstable oscillations.

PPO avoids these structural issues because:

- It does not use a max over Q-values
- It performs on-policy updates
- Value estimates are auxiliary rather than directly driving action selection

6.3 PPO Clipping and Generalised Advantage Estimation (GAE)

PPO Clipped Objective

PPO constrains policy updates through the clipped surrogate objective:

$$L^{CLIP}(\theta) = \mathbb{E}[\min(r_t(\theta)A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)]$$

where:

$$r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$$

This objective:

- Approximates a trust-region method
- Prevents excessively large updates
- Improves training stability

Generalised Advantage Estimation (GAE)

The advantage function is defined as:

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

GAE computes:

$$A_t^{GAE(\lambda)} = \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l}$$

where the temporal-difference residual is:

$$\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$$

The parameter $\lambda \in [0,1]$ controls the bias-variance trade-off:

- $\lambda = 0$: lower variance, higher bias
- $\lambda = 1$: lower bias, higher variance

GAE reduces variance in policy gradient updates while maintaining acceptable bias, improving sample efficiency and stability.

6.4 Epsilon-Greedy vs Entropy-Driven Exploration

Epsilon-Greedy Exploration (DQN)

The behaviour policy is defined as:

$$\pi(a|s) = \begin{cases} \text{random action} & \text{with probability } \epsilon \\ \underset{a}{\operatorname{argmax}} Q(s, a) & \text{otherwise} \end{cases}$$

Limitations:

- Uniform and uninformed exploration
- Independent of uncertainty
- Becomes fully greedy as $\epsilon \rightarrow 0$

Exploration is externally imposed rather than integrated into optimisation.

Entropy-Regularised Exploration (PPO)

PPO includes an entropy bonus in the objective:

$$L = L^{CLIP} + \beta \mathbb{E}[H(\pi_\theta(\cdot | s))]$$

where entropy is defined as:

$$H(\pi) = - \sum_a \pi(a|s) \log \pi(a|s)$$

This mechanism:

- Encourages stochastic policies early in training
- Gradually reduces randomness
- Couples exploration with optimisation

Unlike epsilon-greedy, exploration is intrinsic to the learning objective.

6.5 Conceptual Summary

In summary, value-based methods optimise an approximation of the value function and derive the policy implicitly through action maximisation. This structural reliance on the max operator and bootstrapping introduces overestimation bias and makes deep value-based methods particularly sensitive to instability when combined with function approximation and off-policy learning. In contrast, policy-based methods directly parameterise and optimise the policy, avoiding explicit maximisation over value estimates and thereby reducing structural sources of instability. PPO further improves stability through its clipped surrogate objective, which constrains policy updates, and through Generalised Advantage Estimation, which manages the bias-variance trade-off in gradient estimates. Regarding exploration, epsilon-greedy strategies impose external and uninformed randomness that vanishes as epsilon decays, whereas entropy-regularised exploration integrates stochasticity directly into the optimisation objective, allowing exploration to adapt naturally as learning progresses. Overall, these theoretical differences explain why PPO typically exhibits smoother and more stable learning dynamics compared to DQN in complex environments.